

MEDS ML4H 2025 Tutorial

Matthew McDermott

Assistant Professor, Columbia University Department of Biomedical Informatics

Preparation: Suhana Bedi, Teya Bergamaschi, Hyewon Jeong, Simon Lee, Nassim Oufattolle, Patrick Rockenschaub, Kamilė Stankevičiūtė, Ethan Steinberg, Jimeng Sun, Robin van de Water, Michael Wornow, John Wu, Zhenbang Wu, and Justin Xu

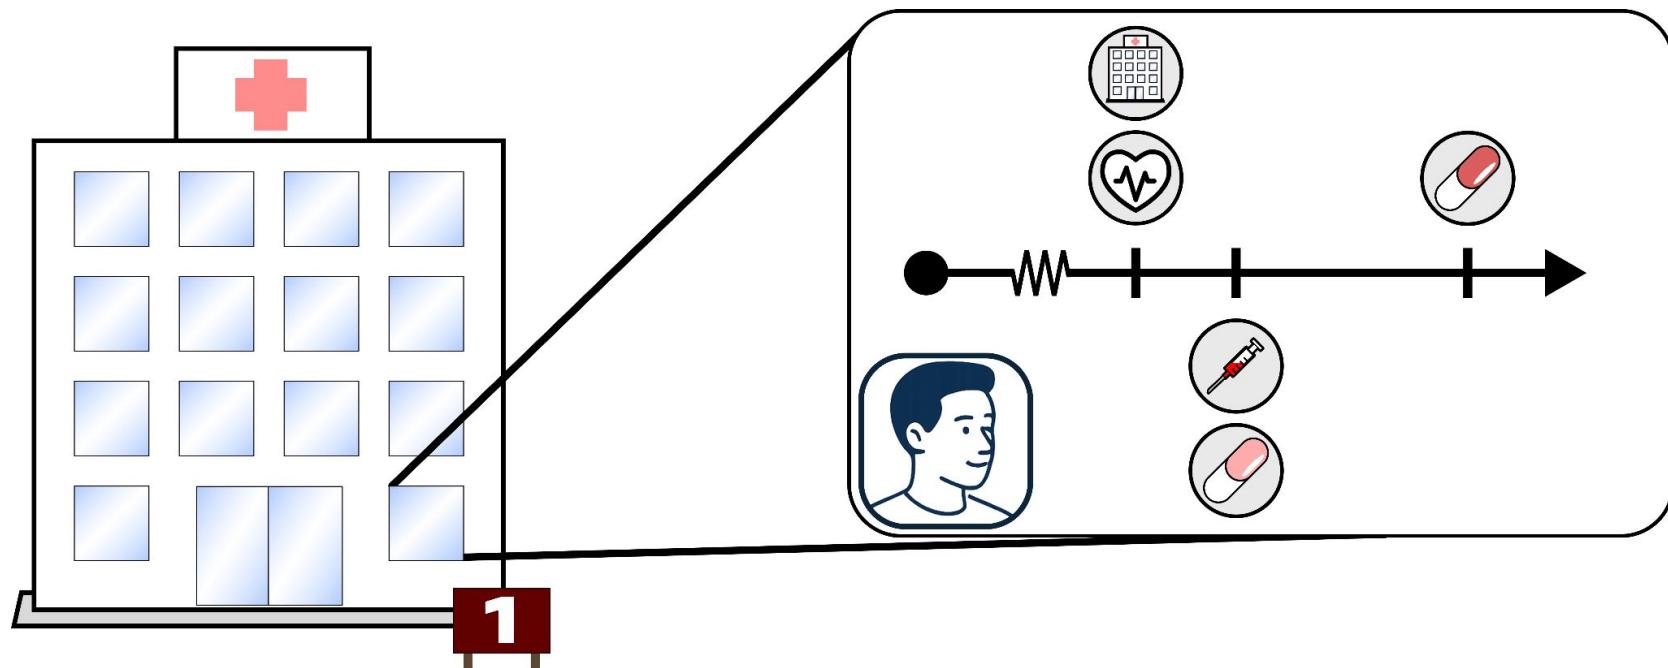


MEDS

<https://medical-event-data-standard.github.io/>

mattmcdermott8@gmail.com

Area: AI for Structured Electronic Health Record (EHR) Data



Learning Goals

1. Understand the core concepts and design principles of the MEDS ecosystem
 - a. *Apply*: Be able to transform data into the MEDS format.
 - b. *Analyze*: Be able to break down a problem into different parts that can leverage different tools.
2. Understand how to build models over MEDS data
 - a. *Understand*: Identify how to map data into the necessary formats for effective ML/AI
3. Identify how to participate in the MEDS community and promote reproducible research
 - a. *Apply*: Use existing models via MEDS-DEV to compare to the state-of-the-art

Outline

- ✓ 9:15 - 9:20 Opening remarks
- CO 1. 9:20 - 9:40 What is MEDS?
- CO 2. 9:40 - 9:45 Guiding Problem Setup
- CO 3. 9:45 - 10:05 Convert your data into MEDS
- CO 4. 10:05 - 10:10 Extract a prediction task cohort with ACES
- CO 5. 10:10 - 10:15 Develop a predictive model 1: Tabular Baseline
- CO 6. 10:15 - 10:35 Develop a predictive model 2: Neural Network
- 7. 10:35 - 10:40 Closing Remarks



[www.tinyurl.com/
MEDS-ML4H-pres](http://www.tinyurl.com/MEDS-ML4H-pres)



Shared Disclaimers

Disclaimer S1: MEDS is a “data-as-interface” ecosystem. This means that, in general, rather than operating through an interface of python objects, MEDS tools operate through an interface of data in a format. This makes the ecosystem very open, enabling many ways of doing things and many tools that could be used for the same job, not just the tools we show here.

Disclaimer S2: Most MEDS tools have been designed by academic volunteers, and it is more than plausible that there are better, faster, or more efficient ways to do things. If you think you can find one -- great! Put it out there in the community, and we'll gladly use it.

Part 1: What is MEDS and why use it?

Driving Problem: Health AI has a Reproducibility Crisis

We reproduced datasets for 38 experiments corresponding to 28 published studies using MIMIC. In half of the experiments, the sample size we acquired was 25% greater or smaller than the sample size reported.

- Alistair Johnson et. al., 2017

<https://proceedings.mlr.press/v68/johnson17a.html>

MLH papers scored even more poorly when it came to code release... with only ~21% of the papers we analyzed releasing their code publicly

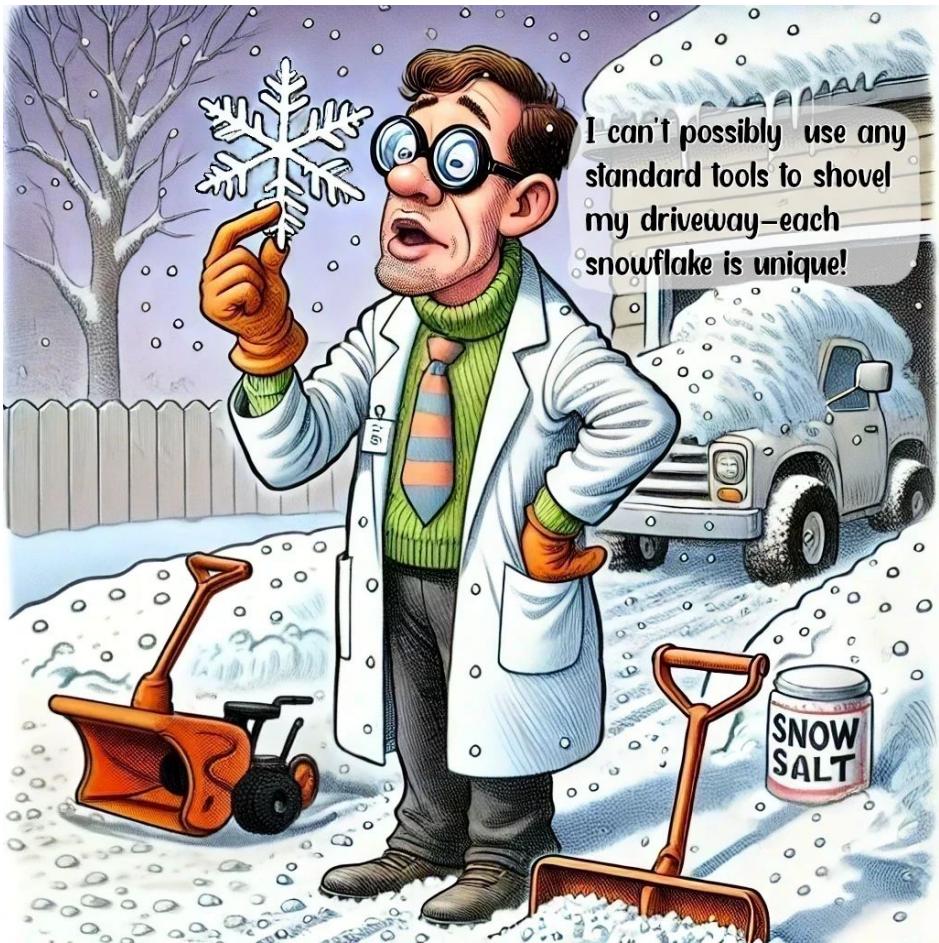
- Matthew McDermott et. al., 2021

<https://doi.org/10.1126/scitranslmed.abb1655>

Of the 218 included articles, 73 (34%) shared code, with 24 (33% of code sharing articles and 11% of all articles) sharing reproducible code

- Kesavan Venkatesh et. al., 2022

<https://doi.org/10.1148/ryai.220081>



The (lack of?) Science of Machine Learning for Healthcare.

McDermott, M. (2025). *Proceedings of the 4th Machine Learning for Health Symposium*, in *Proceedings of Machine Learning Research* 259:19-29

Available from

<https://proceedings.mlr.press/v259/mcdermott25a.html>



To solve this reproducibility crisis, we need a data standard designed for health AI.

Design Principles for a Standard for Health AI

1. Transportability of model training algorithms

The fundamental goal of MEDS is to enable *frictionless* transportability of algorithms first and foremost.

2. Capture the simplest view of the fundamental structure of health data

MEDS asserts that EHR data can be most simply and fundamentally represented by capturing the longitudinal timeline of patient observations. Emphasizing simplicity in this way empowers and simplifies downstream use.

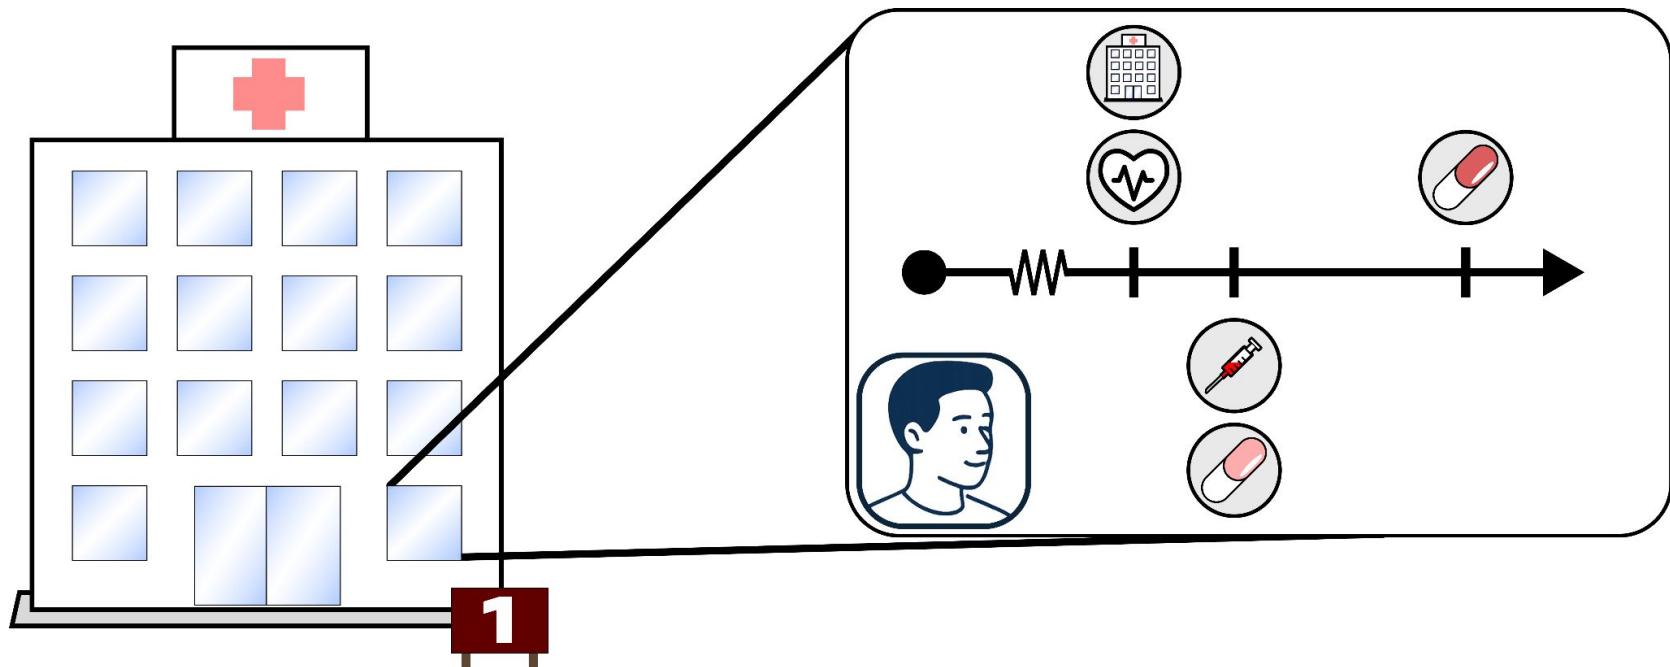
3. Empower development of an open-source ecosystem for health AI

MEDS must enable development of an ecosystem of models, tools, & datasets.

4. Support Computational Performance for Foundation Models

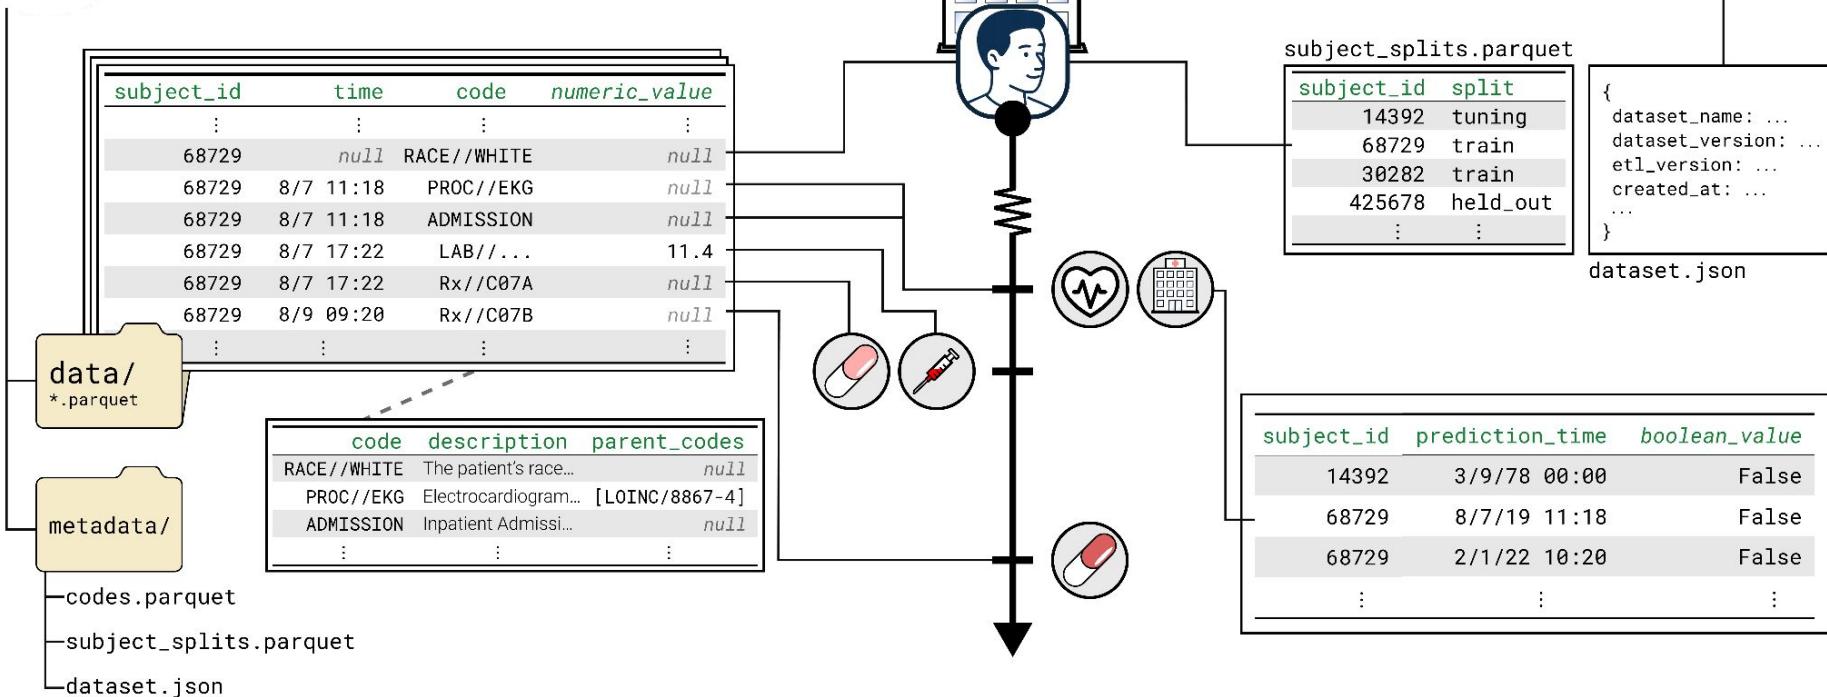
MEDS should support foundation-model scale systems.

A Patient's Data can be represented as a Timeline



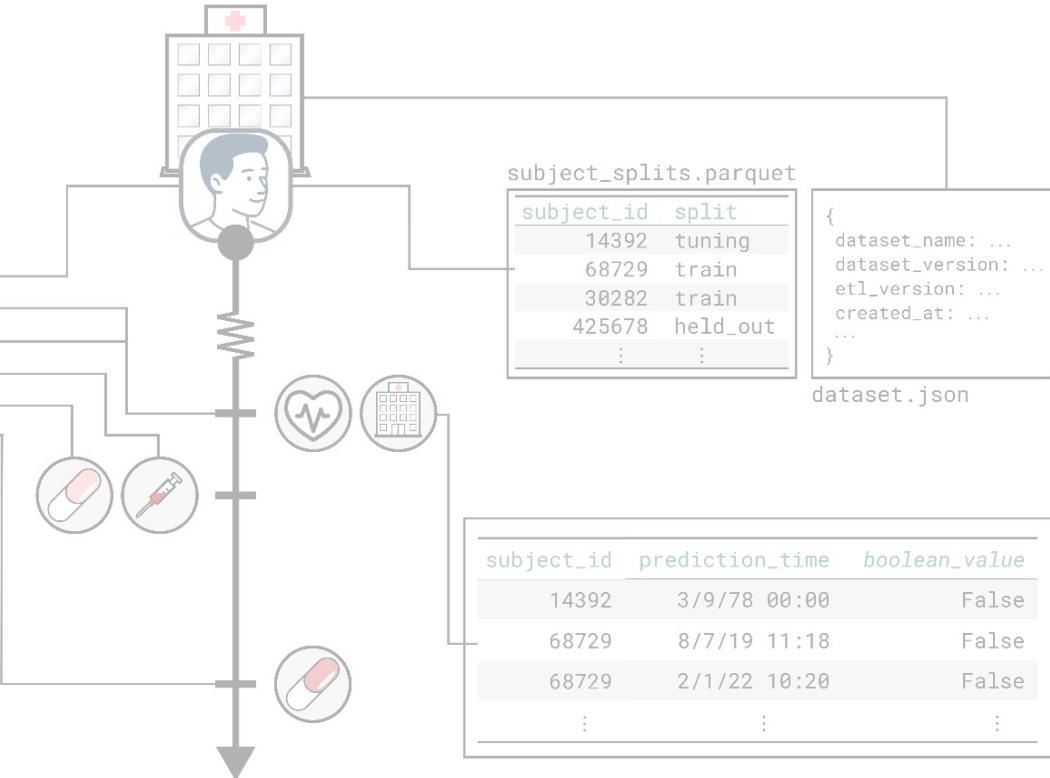
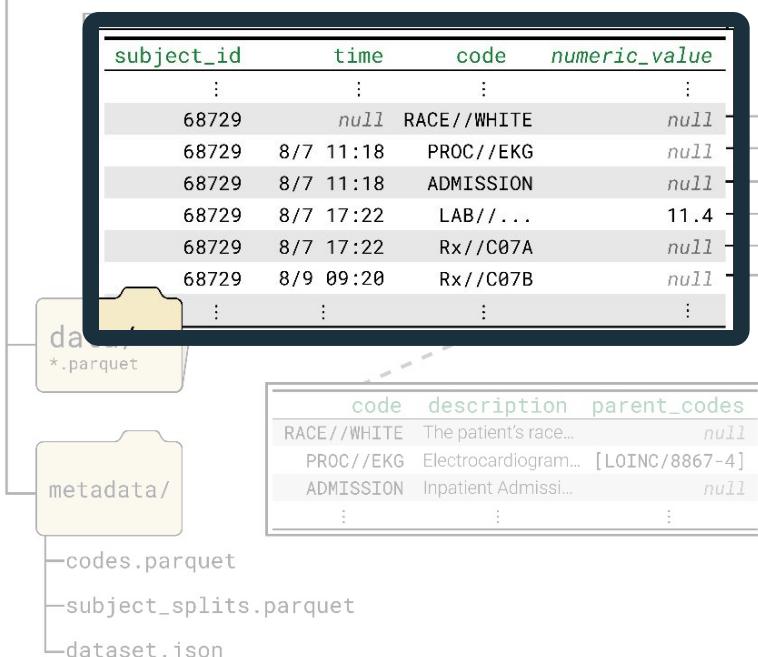


MEDS





MEDS





Data Shards

subject_id	time	code	numeric_value
:	:	:	:
68729	null	RACE//WHITE	null
68729	8/7 11:18	PROC//EKG	null
68729	8/7 11:18	ADMISSION	null
68729	8/7 17:22	LAB//...	11.4
68729	8/7 17:22	Rx//C07A	null
68729	8/9 09:20	Rx//C07B	null
:	:	:	:

- Each shard contains all data for a patient, in time order.



Data Shards

subject_id	time	code	numeric_value
:	:	:	:
68729	null	RACE//WHITE	null
68729	8/7 11:18	PROC//EKG	null
68729	8/7 11:18	ADMISSION	null
68729	8/7 17:22	LAB//...	11.4
68729	8/7 17:22	Rx//C07A	null
68729	8/9 09:20	Rx//C07B	null
:	:	:	:

- Each shard contains all data for a patient, in time order.
- The code column is *unconstrained*.



Data Shards

subject_id	time	code	numeric_value
:	:	:	:
68729	null	RACE//WHITE	null
68729	8/7 11:18	PROC//EKG	null
68729	8/7 11:18	ADMISSION	null
68729	8/7 17:22	LAB//...	11.4
68729	8/7 17:22	Rx//C07A	null
68729	8/9 09:20	Rx//C07B	null
:	:	:	:

- Each shard contains all data for a patient, in time order.
- The code column is *unconstrained*.
- Additional columns can be included as needed.

MEDS data properties and examples

subject_id	time	code	numeric_value
68729		EYE_COLOR//HAZEL	
68729		HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	MEDS_BIRTH	
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	
68729	05/26/2010, 2:30:56	HR	86
68729	05/26/2010, 2:30:56	TEMP	97.8
68729	05/26/2010, 4:51:52	DISCHARGE	
239684		EYE_COLOR//BROWN	
239684		HEIGHT	175.2711152
239684	12/28/1980, 0:00:00	MEDS_BIRTH	
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC	
239684	05/11/2010, 17:41:51	HR	102.6
239684	05/11/2010, 17:41:51	TEMP	96
239684	05/11/2010, 17:48:48	HR	105.1
239684	05/11/2010, 17:48:48	TEMP	96.2

Static observations have null timestamps

subject_id	time	code	numeric_value
68729		EYE_COLOR//HAZEL	160.3953106
68729		HEIGHT	
68729	03/09/1978, 0:00:00	MEDS_BIRTH	
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	
68729	05/26/2010, 2:30:56	HR	86
68729	05/26/2010, 2:30:56	TEMP	97.8
68729	05/26/2010, 4:51:52	DISCHARGE	
239684		EYE_COLOR//BROWN	175.2711152
239684		HEIGHT	
239684	12/28/1980, 0:00:00	MEDS_BIRTH	
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC	
239684	05/11/2010, 17:41:51	HR	102.6
239684	05/11/2010, 17:41:51	TEMP	96
239684	05/11/2010, 17:48:48	HR	105.1
239684	05/11/2010, 17:48:48	TEMP	96.2

*Static
observations*

Observations at the same time share timestamps

*Timestamps
will be
duplicated!*

subject_id	time	code	numeric_value
68729		EYE_COLOR//HAZEL	
68729		HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	MEDS_BIRTH	
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	
68729	05/26/2010, 2:30:56	HR	86
68729	05/26/2010, 2:30:56	TEMP	97.8
68729	05/26/2010, 4:51:52	DISCHARGE	
239684		EYE_COLOR//BROWN	
239684		HEIGHT	175.2711152
239684	12/28/1980, 0:00:00	MEDS_BIRTH	
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC	
239684	05/11/2010, 17:41:51	HR	102.6
239684	05/11/2010, 17:41:51	TEMP	96
239684	05/11/2010, 17:48:48	HR	105.1
239684	05/11/2010, 17:48:48	TEMP	96.2

MEDS mandates standardization, not harmonization

subject_id	time	code	numeric_value	Non-standard vocabularies
68729		EYE_COLOR//HAZEL		
68729		HEIGHT	160.3953106	
68729	03/09/1978, 0:00:00	MEDS_BIRTH		
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY		
68729	05/26/2010, 2:30:56	HR	86	
68729	05/26/2010, 2:30:56	TEMP	97.8	
68729	05/26/2010, 4:51:52	DISCHARGE		
239684		EYE_COLOR//BROWN		
239684		HEIGHT	175.2711152	
239684	12/28/1980, 0:00:00	MEDS_BIRTH		
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC		
239684	05/11/2010, 17:41:51	HR	102.6	
239684	05/11/2010, 17:41:51	TEMP	96	
239684	05/11/2010, 17:48:48	HR	105.1	
239684	05/11/2010, 17:48:48	TEMP	96.2	

MEDS mandates *standardization*, not *harmonization*

subject_id	time	code	numeric_value
68729		EYE_COLOR//HAZEL	
68729		HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	MEDS_BIRTH	

*Non-standard
vocabularies*

Mandate Standardization while remaining agnostic to Harmonization
We can support model training recipe transportability and health AI development without requiring concepts to be aligned with a particular vocabulary.

239684	12/28/1980, 0:00:00	MEDS_BIRTH	
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC	
239684	05/11/2010, 17:41:51	HR	102.6
239684	05/11/2010, 17:41:51	TEMP	96
239684	05/11/2010, 17:48:48	HR	105.1
239684	05/11/2010, 17:48:48	TEMP	96.2

MEDS: minimal standard, no harmonization

subject_id	time	code	numeric_value	Intervals via start...
68729		EYE_COLOR//HAZEL		
68729		HEIGHT	160.3953106	
68729	03/09/1978, 0:00:00	MEDS_BIRTH		
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY		
68729	05/26/2010, 2:30:56	HR	86	
68729	05/26/2010, 2:30:56	TEMP	97.8	
68729	05/26/2010, 4:51:52	DISCHARGE		
239684		EYE_COLOR//BROWN		
239684		HEIGHT	175.2711152	
239684	12/28/1980, 0:00:00	MEDS_BIRTH		
239684	05/11/2010, 17:41:51	ADMISSION//CARDIAC		
239684	05/11/2010, 17:41:51	HR	102.6	
239684	05/11/2010, 17:41:51	TEMP	96	
239684	05/11/2010, 17:48:48	HR	105.1	
239684	05/11/2010, 17:48:48	TEMP	96.2	

... and end events

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.
3. Guarantees of harmonization to external ontologies.

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.
3. Guarantees of harmonization to external ontologies.
4. Canonical representations of different sites for multi-site datasets.

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.
3. Guarantees of harmonization to external ontologies.
4. Canonical representations of different sites for multi-site datasets.
5. Canonical representations of providers for provider-ID augmented datasets

What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.
3. Guarantees of harmonization to external ontologies.
4. Canonical representations of different sites for multi-site datasets.
5. Canonical representations of providers for provider-ID augmented datasets.

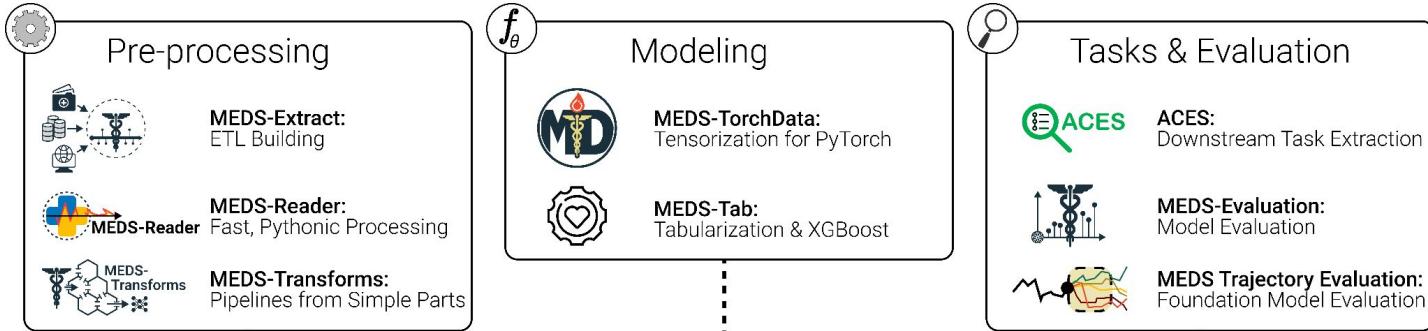
What is MEDS missing?

1. Additional data modalities (e.g., waveforms, images).
 - a. Images are relatively easy to support if needed.
 - b. Waveforms are more challenging as they don't conform to the core MEDS standard in general.
2. Complex representations of "interval" events - instead, separate "start" and "end" events are used.
3. Guarantees of harmonization to external ontologies.
4. Canonical representations of different sites for multi-site datasets.
5. Canonical representations of providers for provider-ID augmented datasets.
6. Natural ways to capture multi-part numeric values (e.g., medication dose, frequency, duration).

... Why should you use it?



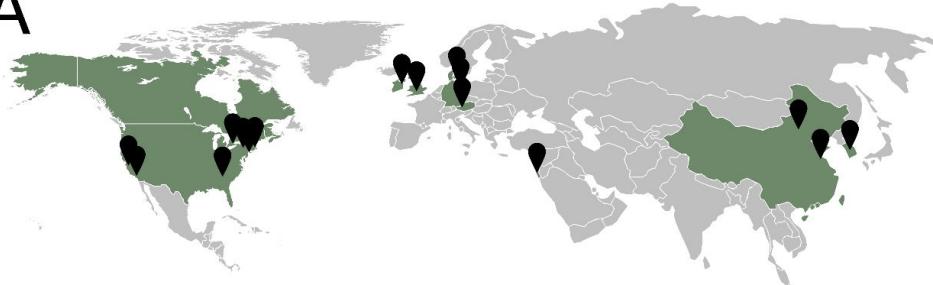
MEDS





MEDS

A



Used in at least 19 institutions

C Up to 100x
Faster



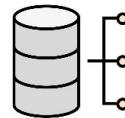
Up to 70% fewer
lines of code



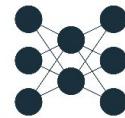
B 14 Papers



17 Datasets



11 Models



More than 13
tools available



Part 1 Colab Notebook: Introduction to MEDS

Time: 5 Minutes

Learning Goals:

1. Ensure the colab notebook setting is familiar.
2. Gain familiarity with the file structure and schema contents of MEDS.



Part 2: Problem Set-up

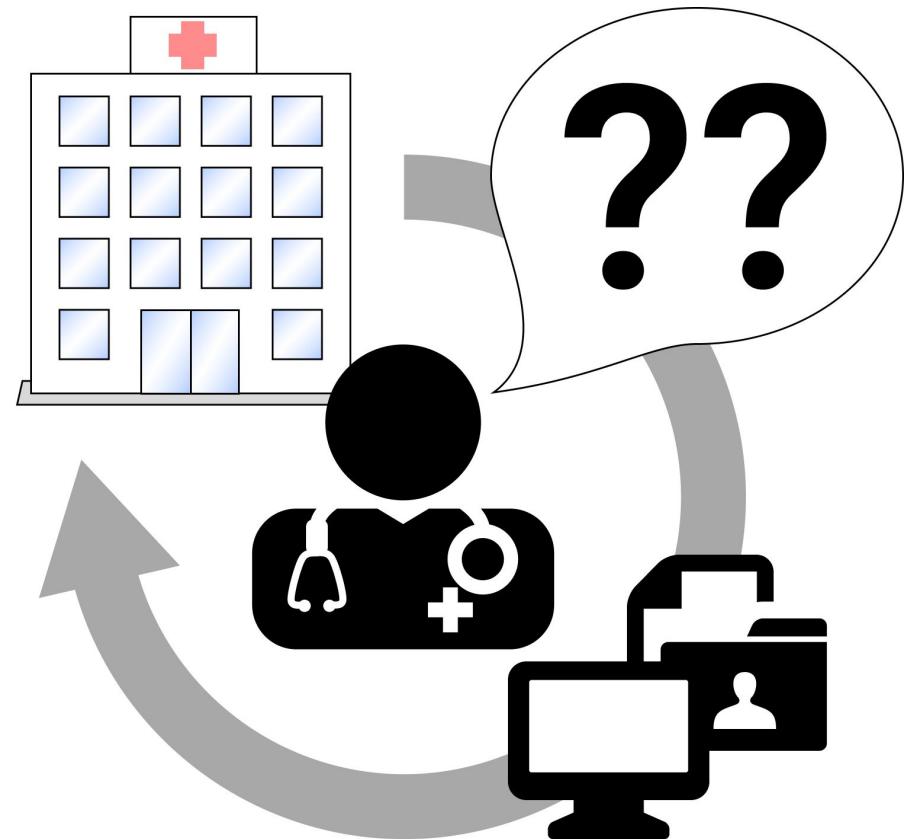
Disclaimers and Goal

This tutorial will be structured around a hypothetical problem. But...

1. The sample problem and setting were chosen for their utility in this tutorial, not for their utility as true clinical tasks!
2. This tutorial will work over *demo data only* -- you should not expect results to be individually meaningful or generalize.
3. Tools and pipelines in this tutorial are configured as they are for educational utility -- do not assume they would necessarily be appropriate in a real modeling task!

Question:

A new colleague asks you: What model should they build for their data?



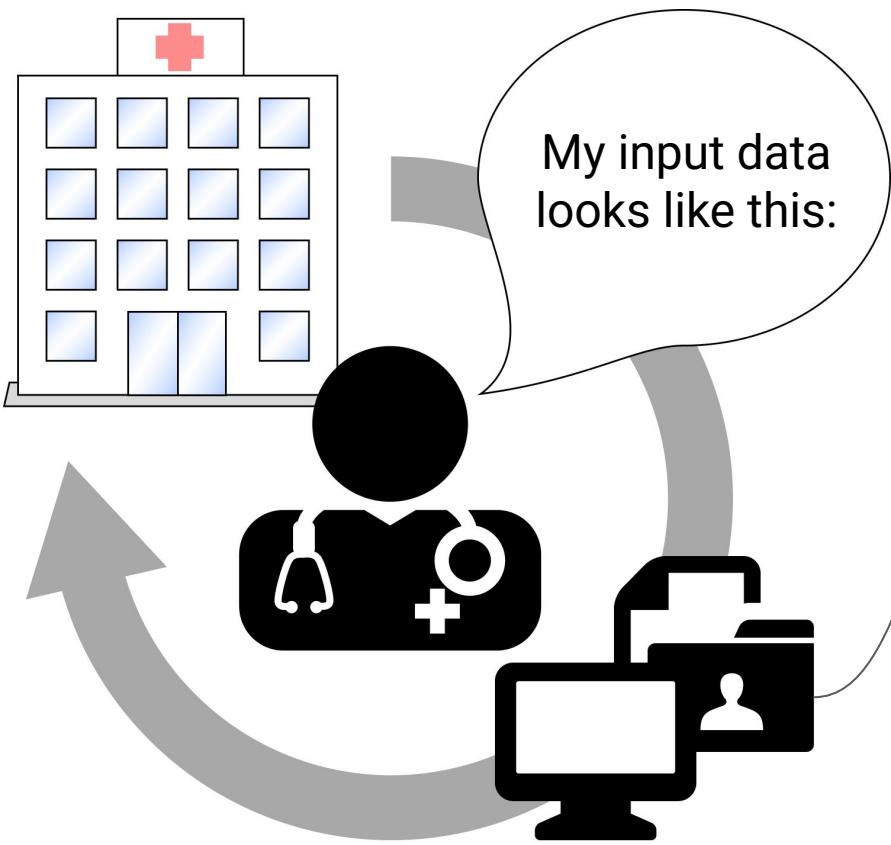
This tutorial

Their data: MIMIC-IV Demo Dataset (<https://physionet.org/content/mimic-iv-demo/2.2/>)

Their problem: *Identify patients who will have a long length of stay in the ICU*

Your goal: *Help them perform this modeling task and identify the right problem framing and model to perform this prediction task, in a computational (i.e., non-deployment) setting*

Part 3: Convert your data into MEDS



```
$ ls $RAW_DATA_DIR
```

hosp/patient.csv.gz
hosp/admissions.csv.gz
icu/charevents.csv.gz

MEDS' Simplicity Enables Easy Extraction Pipelines

```
$ ls $RAW_DATA_DIR
```

hosp/patient.csv

hosp/admissions.csv

icu/chartevents.csv

```
subject_id_col: subject_id
hosp/admissions:
  ed_registration:
    code: ED_REGISTRATION
    time: col(edregtime)
    time_format: "%Y-%m-%d %H:%M:%S"
  admission:
    code:
      - HOSPITAL_ADMISSION
      - col(admission_type)
      - col(admission_location)
    time: col(admittime)
    time_format: "%Y-%m-%d %H:%M:%S"
  insurance: insurance
  language: language
  marital_status: marital_status
  race: race
  hadm_id: hadm_id
...
icu/chartevents:
  event:
    code:
      - LAB
      - col(itemid)
      - col(value uom)
    time: col(charttime)
    time_format: "%Y-%m-%d %H:%M:%S"
    numeric_value: valuenum
    text_value: value
    hadm_id: hadm_id
    icustay_id: stay_id
    _metadata:
      d_labitems_to_loinc:
        description: ["omop_concept_name", "label"]
        itemid: "itemid ({omop_source_code})"
        parent_codes: "{omop_vocabulary_id}/{omop_concept_code}"
        value uom: "value uom"
```

Extracting Data to MEDS with MEDS-Extract

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv

Assumption:

Each row of each file corresponds to one or more *complete* measurements in the output MEDS data view.

Extracting Data to MEDS with MEDS-Extract

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv

Assumption:

Each row of each file corresponds to one or more *complete* measurements in the output MEDS data view.

May require some “*pre-MEDS*” work to validate.

Every row of a MEDS data shard maps to a given input file:

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv

subject_id	time	code	numeric_value
68729	null	EYE_COLOR//HAZEL	null
68729	null	HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	DOB	null
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	null
68729	05/26/2010, 2:30:56	HR	86

Every row of a MEDS data shard maps to a given input file:

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv

subject_id	time	code	numeric_value
68729	<i>null</i>	EYE_COLOR//HAZEL	<i>null</i>
68729	<i>null</i>	HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	DOB	<i>null</i>
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	<i>null</i>
68729	05/26/2010, 2:30:56	HR	86

Every row of a MEDS data shard maps to a given input file:

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv

subject_id	time	code	numeric_value
68729	null	EYE_COLOR//HAZEL	null
68729	null	HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	DOB	null
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	null
68729	05/26/2010, 2:30:56	HR	86

Every row of a MEDS data shard maps to a given input file:

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv

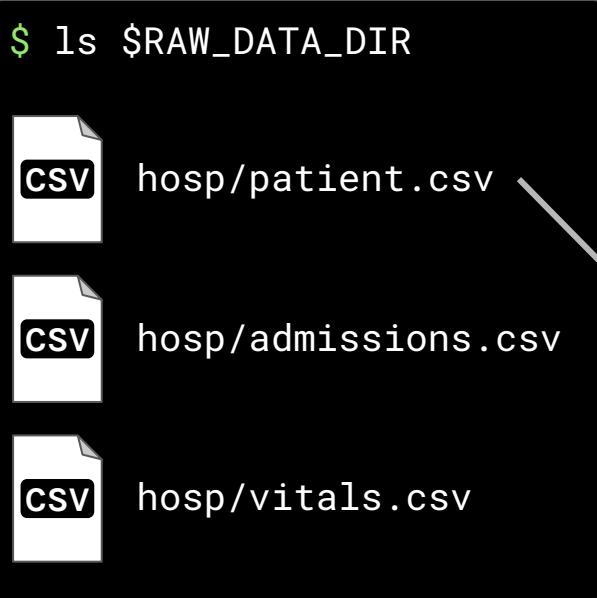


hosp/vitals.csv

subject_id	time	code	numeric_value
68729	null	EYE_COLOR//HAZEL	null
68729	null	HEIGHT	160.3953106
68729	03/09/1978, 0:00:00	DOB	null
68729	05/26/2010, 2:30:56	ADMISSION//PULMONARY	null
68729	05/26/2010, 2:30:56	HR	86

Reversed: Every row of each input file maps to one or more MEDS rows

```
$ ls $RAW_DATA_DIR
```



```
hosp/patient.csv
hosp/admissions.csv
hosp/vitals.csv
```



subject_id	gender	anchor_age	anchor_year	anchor_year_group	dod
100	F	27	2153	2010-2015	null
101	M	54	2124	2015-2020	2144
102	M	78	2155	2010-2015	2160
...

Reversed: Every row of each input file maps to one or more MEDS rows

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv



subject_id	gender	anchor_age	anchor_year	anchor_year_group	dod
100	F	27	2153	2010-2015	null
101	M	54	2124	2015-2020	2144
102	M	78	2155	2010-2015	2160
...

Reversed: Every row of each input file maps to one or more MEDS rows

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv



subject_id	gender	anchor_age	anchor_year	anchor_year_group	dod
100	F	27	2153	2010-2015	null
101	M	54	2124	2015-2020	2144
102	M	78	2155	2010-2015	2160
...

Reversed: Every row of each input file maps to one or more MEDS rows

```
$ ls $RAW_DATA_DIR
```



hosp/patient.csv



hosp/admissions.csv



hosp/vitals.csv



subject_id	gender	anchor_age	anchor_year	anchor_year_group	dod
100	F	27	2153	2010-2015	null
101	M	54	2124	2015-2020	2144
102	M	78	2155	2010-2015	2160
...

MEDS: Extract your data by asking “who”, “what”, and “when”

1. You can think about MEDS data extraction as a repeated task of asking: “To whom is this happening?”, “What is happening?”, and “When is it happening?”
2. This maps to the MEDS-Extract Specification Syntax YAML (MESSY) format: allowing you to communicate and define your extraction parameters.
3. *Bonus:* You can also extract MEDS datasets using other tools or custom pipelines -- only the MEDS output matters!
4. *Bonus:* Lots of datasets already have publicly available MEDS ETLs! Check out this link for more information:



<https://tinyurl.com/MEDS-Datasets>

Part 3 Colab Notebook: Convert to MEDS

Time: 40 Minutes

Learning Goals:

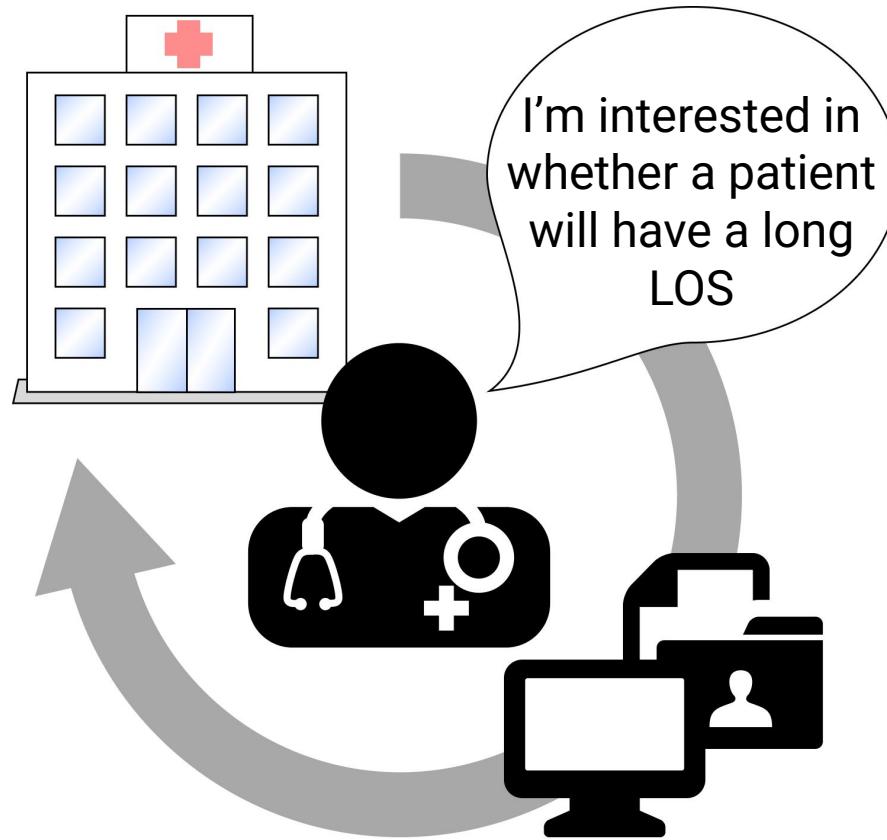
1. Identify MEDS observations in real data.
2. Create a specification using MEDS-Extract that realizes those MEDS observations.

Disclaimers:

1. S1 & S2: Many ways, many tools, and all can be improved!
2. MEDS conventions are not set in stone -- if you participate in the MEDS community, you can help define new conventions for how to do things properly!



Part 4: Identify your prediction task





MEDS Label Schema

The MEDS label schema requires an index (a `subject_id`, and `prediction_time`) and permits optional labels of type including `boolean_value`, `integer_value`, `float_value`, and `categorical_value`. These labels can be predicted using any of the data of the indexed subject that occurred anytime at or before the indexed prediction time.

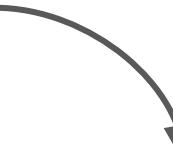
<code>subject_id</code>	<code>prediction_time</code>	<code>boolean_value</code>
68729	3/9/78 00:00	False
68729	5/2/10 14:22	False
68729	5/2/10 14:34	False
125829	4/9/18 18:19	True



MEDS Label Schema

Index DataFrame

The MEDS label schema requires an index (a `subject_id`, and `prediction_time`) and permits optional labels of type including `boolean_value`, `integer_value`, `float_value`, and `categorical_value`. These labels can be predicted using any of the data of the indexed subject that occurred anytime at or before the indexed prediction time.



subject_id	prediction_time	boolean_value
68729	3/9/78 00:00	False
68729	5/2/10 14:22	False
68729	5/2/10 14:34	False
125829	4/9/18 18:19	True



MEDS Label Schema

“Task” or “Label” DataFrame

The MEDS label schema requires an index (a `subject_id`, and `prediction_time`) and permits optional labels of type including `boolean_value`, `integer_value`, `float_value`, and `categorical_value`. These labels can be predicted using any of the data of the indexed subject that occurred anytime at or before the indexed prediction time.



subject_id	prediction_time	boolean_value
68729	3/9/78 00:00	False
68729	5/2/10 14:22	False
68729	5/2/10 14:34	False
125829	4/9/18 18:19	True



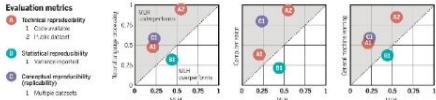
ACES: Reproducible Extraction of Task Cohorts for EHRs

Justin Xu, Jack Gallifant, Alistair E. W. Johnson, Matthew B. A. McDermott



Health AI has a Reproducibility Crisis

Health AI faces a **systemic reproducibility crisis**, limiting our ability to do effective science. We need to build a health AI ecosystem to change that, and ACES builds on that foundation.



ACES Leverages Event-Stream Schemas

The MEDS data schema epitomizes simplicity and has only 4 required columns: subject_id, time, code, and numeric_value.

subject_id	time	code	numeric_value
68720	null	RACE/WHITE	null
68720	SCW/M	SEX/M	null
68720	3/9/78 08:00	MEDS_BIRTH	null
68720	5/2/10 14:22	ED/REG	null
68720	5/2/10 14:34	HR/bpm	93.0
68720	5/2/10 20:00	ED/OUT	null
125829	null	SEX/F	null
125829	4/9/18 18:19	ADMISSION/CARDIAC	null

code	description	parent_codes
RACE/WHITE	The patient's race...	null
SCW/M	The patient's sex...	null
SEX/M	The patient's sex...	null
SEX/F	The patient's sex...	null
MEDS_BIRTH	null	null
HR/bpm	Heart rate, measured...	[L01NC/8867-4]
ED/REG	Emergency department...	null
ED/OUT	Emergency department...	null
SEX/F	The patient's sex...	null

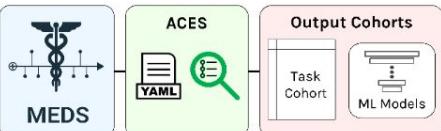
The MEDS codes schema may contain descriptions and links to external ontologies for elements of the code vocabulary.

- Single stream of events!
- Simple and flexible to use!
- Easy transformations!

Learn More about MEDS:



Task Extraction Made Easy!



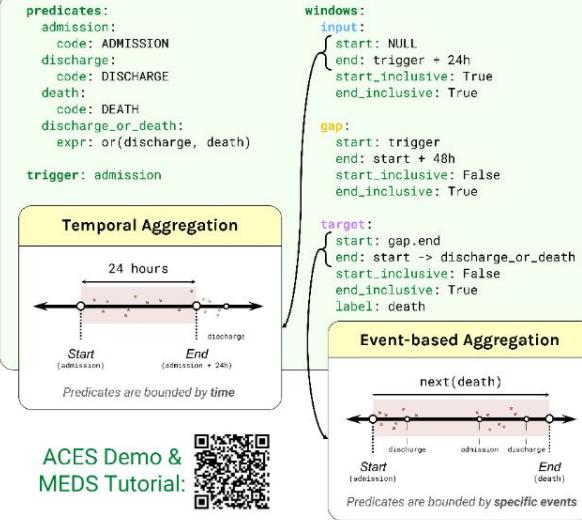
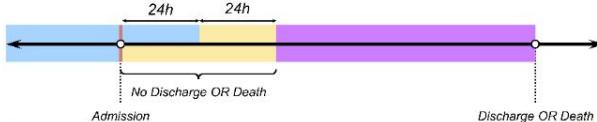
- Transparent
- Reproducible
- Extractable on diverse datasets

ACES Demo & MEDS Tutorial:

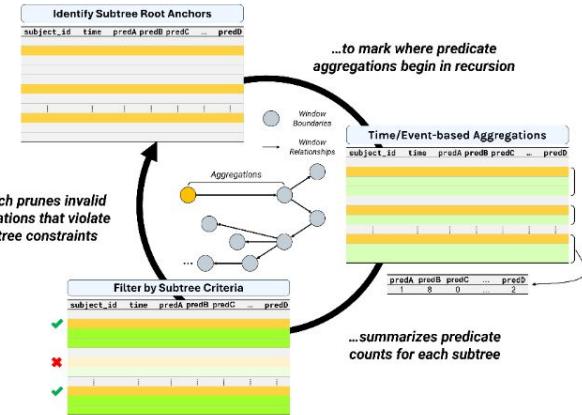


ACES Configuration Files

In-hospital Mortality: Given the first 24 hours of a patient's stay, predict whether or not they will die within this hospital admission, with a gap time of 48h.



Recursive Algorithm for Extraction



How can ACES help you?

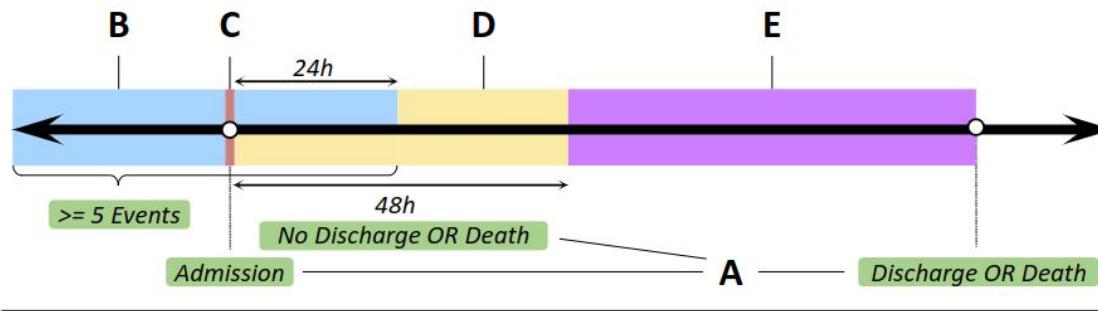
- For a variety of pre-defined tasks, check out (or contribute to) MEDS-DEV:
- For more info on how to write your own ACES configs, check out our documentation:



Acknowledgements

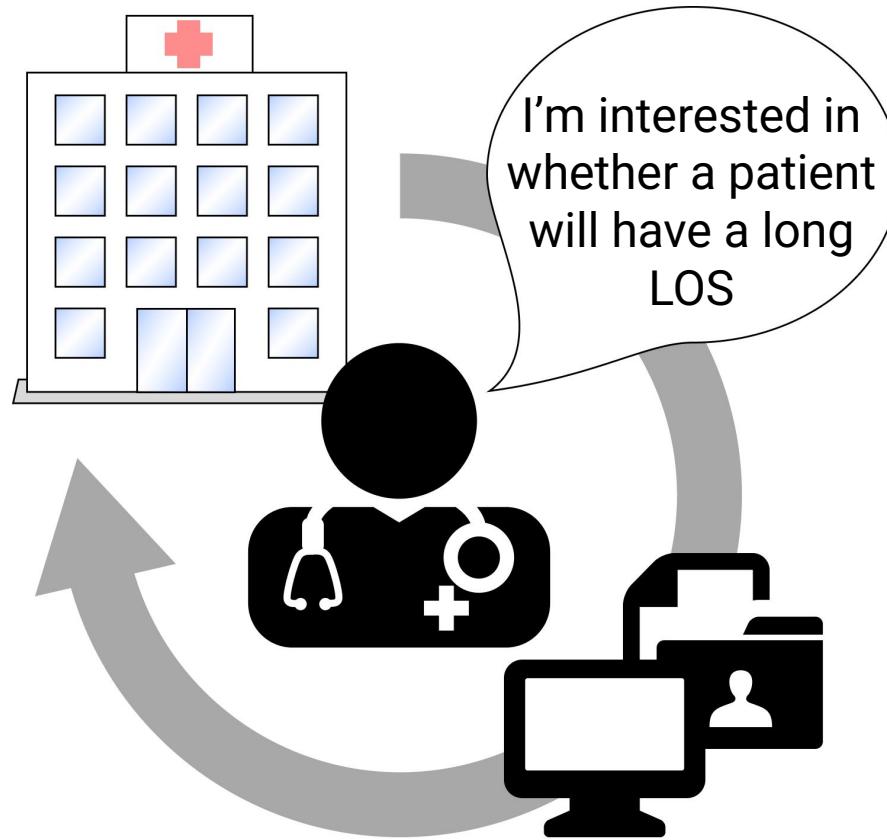
MBAM gratefully acknowledges support from a Berkowitz Postdoctoral Fellowship at Harvard Medical School. JG is funded by the National Institutes of Health through NIH-NIA R01CA294033. JG greatly appreciates support from supervisors David Tyre (University of Oxford) and Curtis Langlotz (Stanford University). We also acknowledge valuable contributions by Tim Pollard (Massachusetts Institute of Technology) and by the broader MEDS ecosystem of contributors and users.

In-hospital Mortality Prediction



Task Configuration

<p>A</p> <p>predicates:</p> <ul style="list-style-type: none"> admission: code: ADMISSION discharge: code: DISCHARGE death: code: DEATH discharge_or_death: expr: or(discharge, death) 	<p>B</p> <p>input:</p> <ul style="list-style-type: none"> start: NULL end: trigger + 24 hours start_inclusive: True end_inclusive: True has: <p>_ANY_EVENT: (5, None)</p>
<p>C</p> <p>trigger: admission</p>	<p>D</p> <p>gap:</p> <ul style="list-style-type: none"> start: trigger end: start + 48 hours start_inclusive: False end_inclusive: True has: <p>admission: (None, 0) discharge: (None, 0) death: (None, 0)</p>
<p>E</p> <p>target:</p> <ul style="list-style-type: none"> start: gap.end end: start -> discharge_or_death start_inclusive: False end_inclusive: True label: death 	



Measure what you value



Tasks can be discussed by the community!

How should we handle censoring / loss-of-follow-up for readmission risk prediction? #9

 Closed  #183

 mmcdermott opened on Aug 14, 2024 · edited by mmcdermott

Edits Member ...

For example, consider the [readmission/general_hospital/30d](#) task.

Tagging for comments: [@prockenschaub](#) [@shalmajoshi](#) [@justin13601](#) [@tompollard](#) [@jwoo5](#)

As of the commit referenced in the link above, this task currently excludes all patients who do not have one data element after 30 days. That may or may not be advisable.

Let's define x to be a patient's data as of a prediction time, R to be the label of whether or not there is an admission event within 30 days ($R = 1$ if so, $R = 0$ otherwise), E to be a binary variable indicating whether or not we have data after 30 days ($E = 1$ if there is data after 30 days, $E = 0$ otherwise), and M to be a binary variable indicating whether or not the patient dies within 30 days ($M = 1$ if they do, $M = 0$ otherwise).

Note that our constraints here are that we are limited to only a binary classification task. We can't, within this current scope, change frameworks to a survival analysis or something.

The question here is whether to include, exclude, and what label to assign for training this task. There are a few options:

Option 1: Only predict for patients with data observed more than 30 days out

Description	R	E	M	Include/Exclude	Training label y
Patients w/ data after 30 days are included & labeled	*	1	0	Include	R
Patients w/o data after 30 days are excluded	*	0	*	Exclude	N/A



Assignees

No one - Assign yourself

Labels



Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

 Code with agent mode

Tasks can be discussed by the community!

How should we handle censoring / loss-of-follow-up for readmission risk prediction? #9

 Closed  #183

 mmcdermott opened on Aug 14, 2024 · edited by mmcdermott

Edits Member ...

For example, consider the [readmission/general_hospital/30d](#) task.

Tagging for comments: [@prockenschaub](#) [@shalmajoshi](#) [@justin13601](#) [@tompollard](#) [@jwoo5](#)

As of the commit referenced in the link above, this task currently excludes all patients who do not have one data element after 30 days. That may or may not be advisable.

Let's define x to be a patient's data as of a prediction time, R to be the label of whether or not there is an admission event within 30 days ($R = 1$ if so, $R = 0$ otherwise), E to be a binary variable indicating whether or not we have data after 30 days ($E = 1$ if there is data after 30 days, $E = 0$ otherwise), and M to be a binary variable indicating whether or not the patient dies within 30 days ($M = 1$ if they do, $M = 0$ otherwise).

Note that our constraints here are that we are limited to only a binary classification task. We can't, within this current scope, change frameworks to a survival analysis or something.

The question here is whether to include, exclude, and what label to assign for training this task. There are a few options:

Option 1: Only predict for patients with data observed more than 30 days out

Description	R	E	M	Include/Exclude	Training label y
Patients w/ data after 30 days are included & labeled	*	1	0	Include	R
Patients w/o data after 30 days are excluded	*	0	*	Exclude	N/A

 Edit  New issue 

 Assignees  Labels  Type  Projects  Milestone  Relationships  Development

No one - Assign yourself

Tasks

help wanted

priority:high

question

No type

No projects

No milestone

None yet

 Code with agent mode



Final task:

Predict using the first 24 hours of data if the ICU stay will be less than 3 days long.

```
● ● ●

predicates:
  icu_admission:
    code: { regex: "^ICU_ADMISSION//.*" }
  icu_discharge:
    code: { regex: "^ICU_DISCHARGE//.*" }
  death:
    code: { regex: "MEDS_DEATH.*" }

  # CMO predicates
  cmo_1:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "Comfort measures only"
  cmo_2:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "Comfort care (CMO, Comfort Measures)"

  # DNR predicates
  dnr_1:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "DNR / UNT"
  dnr_2:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "DNR (Do Not Attempt Resuscitation) [DNR]"
  dnr_3:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "DNR (Do Not Attempt Resuscitation) [DNR] / DNI"
  dnr_4:
    code: { any: ["LAB//220001//UNK", "LAB//223758//UNK"] }
    text_value: "DNR (do not resuscitate)"

  # derived predicates
  cmo:
    expr: or(cmo_1, cmo_2)
  dnr:
    expr: or(dnr_1, dnr_2, dnr_3, dnr_4)

trigger: icu_admission

windows:
  input:
    start: null
    end: trigger + 24h
    start_inclusive: True
    end_inclusive: True
    index_timestamp: end
    has:
      cmo: (None, 0) # Exclude patients on comfort measures only
      dnr: (None, 0) # Exclude patients with DNR orders
  gap:
    start: trigger
    end: start + 30h
    start_inclusive: False
    end_inclusive: True
    has:
      cmo: (None, 0)
      dnr: (None, 0)
      icu_discharge: (None, 0)
target:
  start: trigger
  end: start + 3d
  start_inclusive: True
  end_inclusive: True
  label: icu_discharge
  has:
    death: (None, 0)
```

Part 4 Colab Notebook: Develop an ACES File

Time: 20 Minutes

Learning Goals:

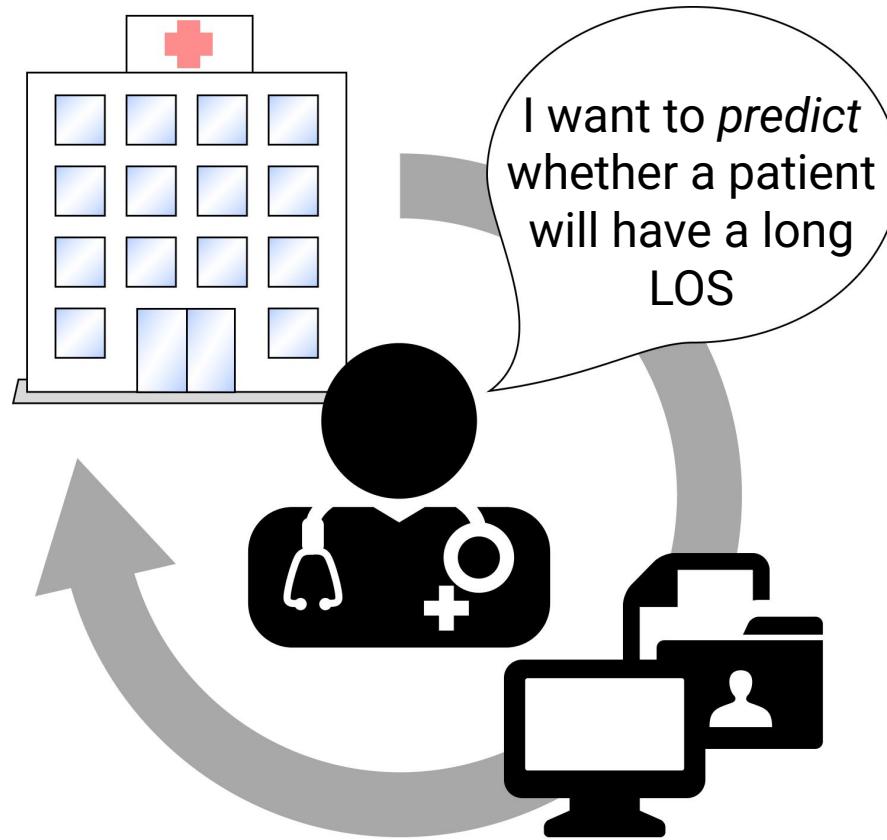
1. *Examine, question, and refine* a stated prediction goal into a meaningful predictive cohort.
2. Create an ACES configuration file to extract that cohort.

Disclaimers:

1. S1 & S2: Many ways, many tools, and all can be improved!
2. This task and sample “conversation” is not necessarily representative of a real discussion with a clinician.



Part 5: Building a tabular baseline model



Part 5 Colab Notebook: Build a Tabular Model

Time: 30 Minutes

Learning Goals:

1. *Explain* tabular baseline modeling needs and how MEDS supports them.
2. *Assemble* a tabular baseline model for the MEDS data using default MEDS tools to make a prediction.

Disclaimers:

1. S1 & S2: Many ways, many tools, and all can be improved!
2. What model will work best on your data? Nobody knows but you (eventually)



Part 5 Takeaways:

1. Manual tabularization is not the primary goal of MEDS, but it is very tractable with sufficient data expertise.
2. Tabularization can be seen as asking “What measurements, over what windows relative to the prediction time, with what aggregation functions?”
3. Tabularization is challenging to do efficiently at large scale.

Part 6: Neural Network Models

stages:

- `fit_normalization`:
- `aggregations`:
 - "code/n_occurrences"
 - "code/n_subjects"
 - "values/n_occurrences"
 - "values/sum"
 - "values/sum_sqd"
- `fit_vocabulary_indices`
- `normalization`
- `tokenization`
- `tensorization`

If your input MEDS dataset lives in `$MEDS_ROOT` and you want to store your pre-processed files in `$PYD_ROOT`, you run:

```
MTD_preprocess MEDS_dataset_dir="$MEDS_ROOT" output_dir="$PYD_ROOT"
```

Step 3: Use the dataset

To use a dataset, you need to (1) define your configuration object and (2) create the dataset object. The only required configuration parameters are `tensorized_cohort_dir`, which points to the root directory containing the pre-processed data on disk (`$PYD_ROOT` in the above example), and `max_seq_len`, which is the maximum sequence length you want to use for your model. Here's an example:

```
import os
from meds_torchdata import MEDSPytorchDataset, MEDSTorchDataConfig

cfg = MEDSTorchDataConfig(tensorized_cohort_dir=os.environ["PYD_ROOT"], max_seq_len=512)
pyd = MEDSPytorchDataset(cfg, split="train")
```

Part 6 Colab Notebook: Build a Neural Network Model

Time: 30 Minutes

Learning Goals:

1. *Explain* longitudinal neural network modeling needs and how MEDS supports them.
2. *Assemble* a neural network models *using* default MEDS tools to make a prediction.

Disclaimers:

1. S1 & S2: Many ways, many tools, and all can be improved!
2. What model will work best on your data? Nobody knows but you (eventually)



Part 6 Takeaways:

1. Neural network models have very different needs than tabular models.
2. For longitudinal neural network models, you need to ask how you pad or truncate sequences, how you sample sequences during training, and how you tokenize and tensorize your data.
3. Efficiency of data loading and processing is highly important.
4. MEDS-TorchData lets you build PyTorch models over MEDS datasets out of the box.

Conclusion

Part 1	Learn about the MEDS ecosystem	Explain why MEDS enables a collaborative, reproducible, open-source ecosystem in Health AI
Part 3	Convert new data into the MEDS format	Understand the conceptual and technical specification of MEDS-Extract
Part 4	Learn how to define prediction tasks with ACES	<ul style="list-style-type: none">• Identify and decompose a task into operationalized targets• Identify necessary predicates and criteria to define a common prediction task.
Part 5	Learn how to build baseline and neural network models using MEDS tools	<ul style="list-style-type: none">• Separate tabularization and baseline model needs from NN needs.• Leverage MEDS ecosystem tools like MEDS-Tab, MEDS Transforms, and MEDS TorchData to accelerate modeling.
Part 6	Learn how to leverage the research of the community to empower your own	

MEDS-DEV:

The MEDS Decentralized, Extensible, Validation Benchmark

If reproducibility is made trivial, we can realize all aspects of assessing a Health AI algorithm under a simple, easy to use interface

```
# Build datasets:  
meds-dev-dataset dataset=$DATASET_NAME output_dir=$DATASET_DIR  
  
# Extract tasks:  
meds-dev-task task=$TASK_NAME dataset=$DATASET_NAME output_dir=$LABELS_DIR dataset_dir=$DATASET_DIR  
  
# Train models:  
# 1. Pre-train a model on unsupervised data  
meds-dev-model model=$MODEL_NAME dataset_dir=$DATASET_DIR mode=train dataset_type=unsupervised  
output_dir=$PRETRAINED_MODEL_DIR  
  
# 2. Fine-tune a model on supervised data  
meds-dev-model model=$MODEL_NAME dataset_dir=$DATASET_DIR labels_dir=$LABELS_DIR mode=train dataset_type=supervised  
output_dir=$FINETUNED_MODEL_DIR model_initialization_dir=$PRETRAINED_MODEL_DIR  
  
# 3. Make predictions with a model for the held-out set:  
meds-dev-model model=$MODEL_NAME dataset_dir=$DATASET_DIR labels_dir=$LABELS_DIR mode=predict dataset_type=supervised  
split=held_out output_dir=$PREDICTIONS_DIR model_initialization_dir=$FINETUNED_MODEL_DIR
```

Acknowledgements



**Kamilė
Stankevičiūtė**
Ph.D. Student, University of
Cambridge



Justin Xu
Ph.D. Student, University of
Oxford



Ethan Steinberg
Researcher, Prealize Health



Acknowledgements

Edward Choi
Ethan Steinberg
Jason A. Fries
Jungwoo Oh
Matthew McDermott

Michael Wornow
Nigam H. Shah
Patrick Rockenschaub
Robin P. van de Water
Tom J. Pollard

Aleksia Kolo
Chao Pang
Edward Choi
Ethan Steinberg
Hyewon Jeong
Jack Gallifant
Jason A. Fries
Jeffrey N. Chiang

Jungwoo Oh
Justin Xu
Kamilė Stankevičiūtė
Kiril V. Klein
Matthew McDermott
Mikkel Odgaard
Nassim Oufattolle
Patrick Rockenschaub
Pawel Renc

Robin van de Water
Shalmali Joshi
Simon A. Lee
Teya S. Bergamaschi
Tom J. Pollard
Vincent Jeanselme
Young Sang Choi



MEDS Publications

MEDS Working Group: Bert Arnrich, Edward Choi, Jason A. Fries, Matthew B. A. McDermott, Jungwoo Oh, Tom J. Pollard, Nigam Shah, Ethan Steinberg, Michael Wornow, Robin van de Water. 2024. "Medical Event Data Standard (MEDS): Facilitating Machine Learning for Health." In ICLR 2024 Workshop TS4H.

[https://openreview.net/forum?id=lsHy2ebjIG&referrer=%5BAuthor%20Console%5D\(%2Fgroup%3Fid%3DICLR.cc%2F2024%2FWorkshop%2FTS4H%2FAuthors%23your-submissions\).](https://openreview.net/forum?id=lsHy2ebjIG&referrer=%5BAuthor%20Console%5D(%2Fgroup%3Fid%3DICLR.cc%2F2024%2FWorkshop%2FTS4H%2FAuthors%23your-submissions).)

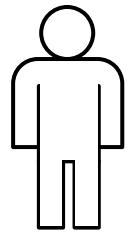
Oufattolle, Nassim, Teya Bergamaschi, Aleksia Kolo, Hyewon Jeong, Hanna Gaggin, Collin M. Stultz, and Matthew B. A. McDermott. 2024. "MEDS-Tab: Automated Tabularization and Baseline Methods for MEDS Datasets." *arXiv [Cs.LG]*. arXiv. <http://arxiv.org/abs/2411.00200>.

Steinberg, E., Michael Wornow, Suhana Bedi, J. Fries, Matthew B. A. McDermott, and Nigam H. Shah. 2024. "Meds_reader: A Fast and Efficient EHR Processing Library." In *Machine Learning for Health Symposium (Findings Track)*. Vol. abs/2409.09095. <https://doi.org/10.48550/arXiv.2409.09095>.

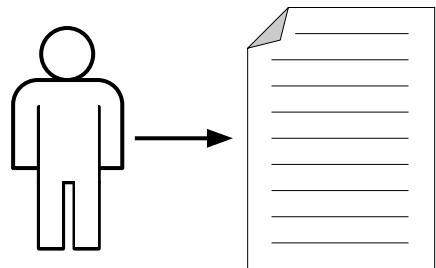
Xu, Justin, Jack Gallifant, Alistair E. W. Johnson, and Matthew B. A. McDermott. 2025. "ACES: Automatic Cohort Extraction System for Event-Stream Datasets." In *Proceedings of the International Conference on Learning Representations (in Press)*.

Reproducibility Crisis

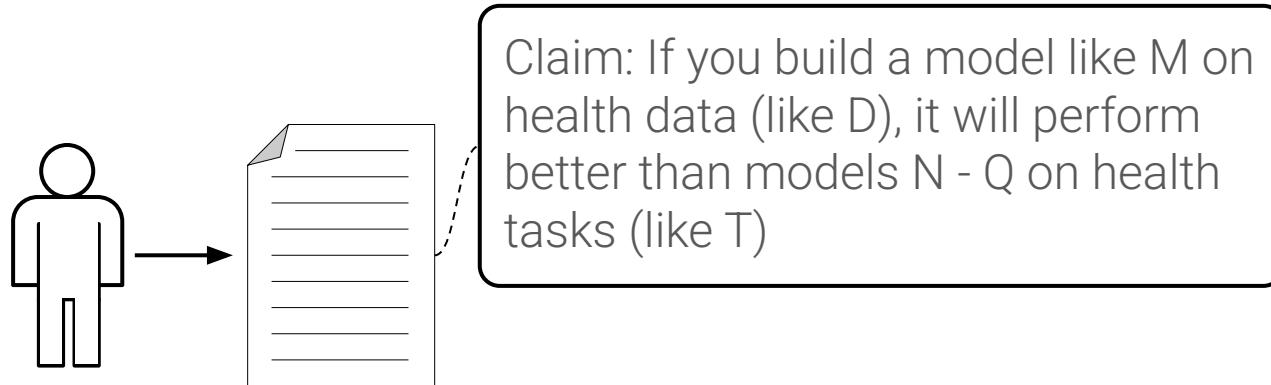
Scientific Process of Health AI



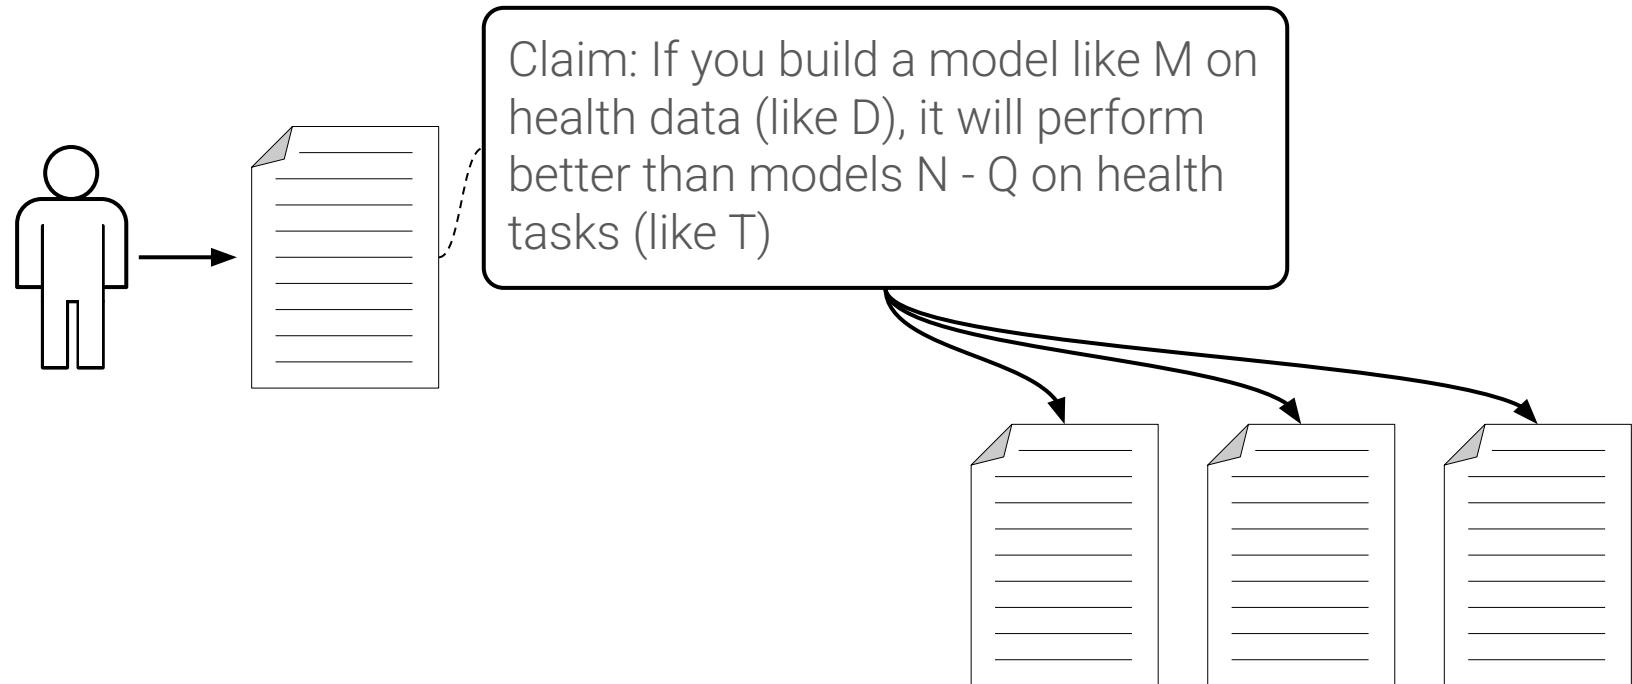
Scientific Process of Health AI



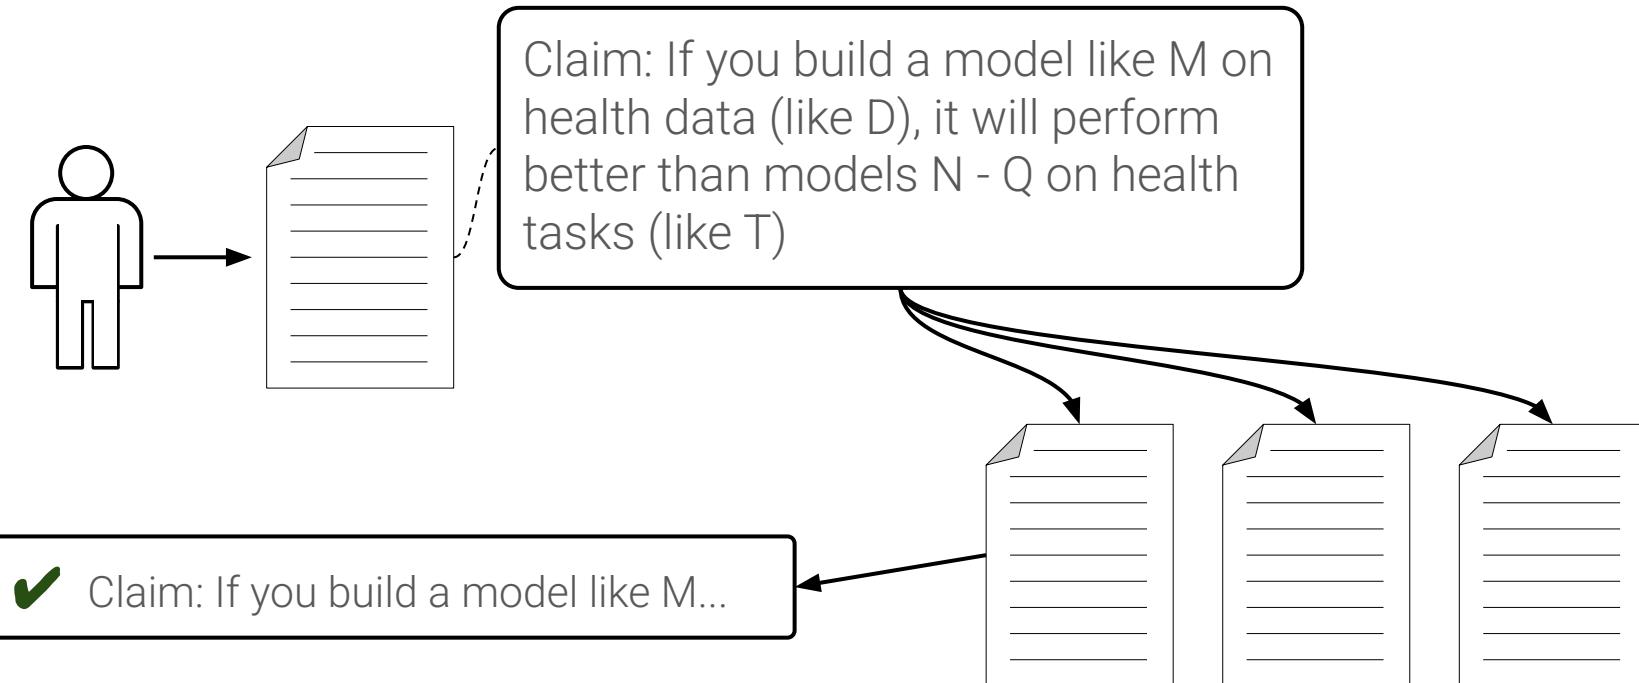
Scientific Process of Health AI: Claims over model training recipe superiority



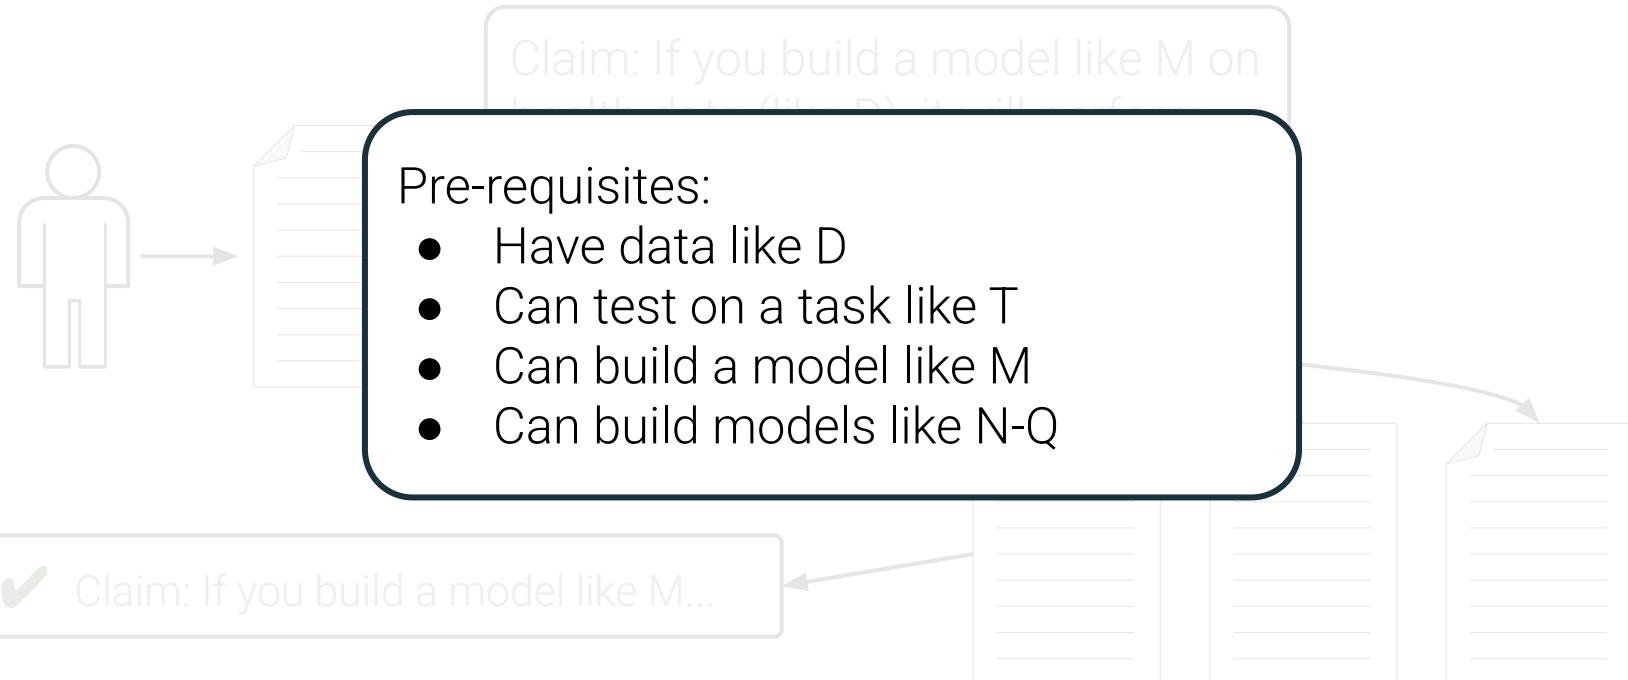
Scientific Process of Health AI: Claims over model training recipe superiority



Scientific Process of Health AI: Reproducibility implies that the same training recipes are superior on new health data



What do we need to check reproducibility?



Assume access to data (or public data)



Pre-requisites:

- ✓ Have data like D
- Can test on a task like T
- Can build a model like M
- Can build models like N-Q

Claim: If you build a model like M on
task like T (M → T) you will get S



Claim: If you build a model like M...



We can't reliably define prediction tasks!

Pre-requisites:

- ✓ Have data like D
- ✗ Can test on a task like T
- Can build a model like M
- Can build models like N-Q

We reproduced datasets for 38 experiments corresponding to 28 published studies using MIMIC. In half of the experiments, the sample size we acquired was 25% greater or smaller than the sample size reported.

- Alistair Johnson et. al., 2017



Claim: If you build a model like M...

We can't reliably re-train models!

Pre-requisites:

- ✓ Have data like D
- ✗ Can test on a task like T
- ✗ Can build a model like M
- ✗ Can build models like N-Q



Claim: If you build a model like M...

Claim: If you build a model like M...
Then it's likely that (like D), it will work on tasks N-Q.

Of the 218 included articles, 73 (34%) shared code, with 24 (33% of code sharing articles and 11% of all articles) sharing reproducible code

- Kesavan Venkatesh et. al., 2022

<https://doi.org/10.1148/ryai.220081>

MLH papers scored even more poorly when it came to code release... with only ~21% of the papers we analyzed releasing their code publicly

- Matthew McDermott et. al., 2021

<https://doi.org/10.1126/scitranslmed.abb1655>

The Health AI “transportability” crisis:
We can’t transport our code, tasks, or
experiments from one paper to
another

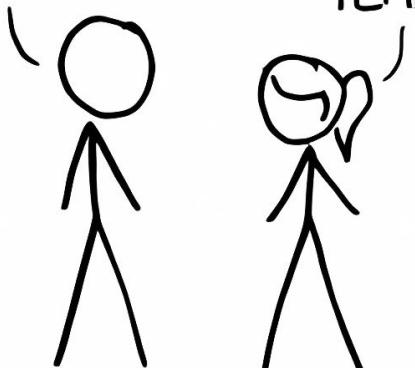
Why not use other
standards?

HOW STANDARDS PROLIFERATE:

(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC)

SITUATION:
THERE ARE
14 COMPETING
STANDARDS.

14?! RIDICULOUS!
WE NEED TO DEVELOP
ONE UNIVERSAL STANDARD
THAT COVERS EVERYONE'S
USE CASES.



YEAH!

SOON:

SITUATION:
THERE ARE
15 COMPETING
STANDARDS.

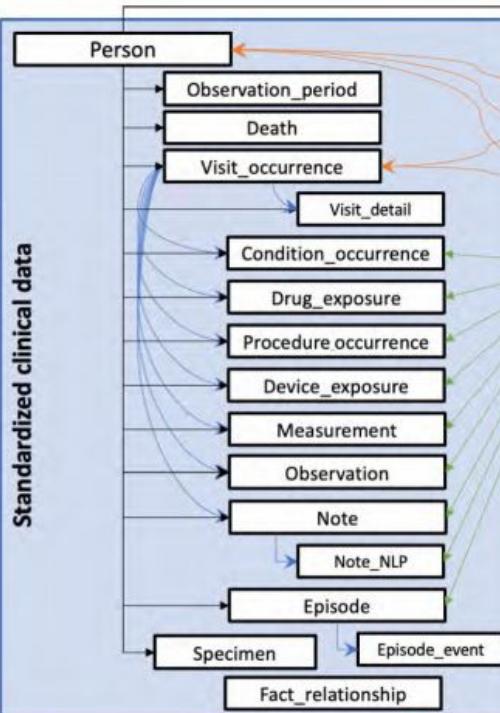
OMOP Common Data Model

The Observational Medical Outcomes Partnership (OMOP) Common Data Model (CDM) is an open community data standard, designed to standardize the structure and content of observational data and to enable efficient analyses that can produce reliable evidence.



"The OMOP Common Data Model serves as the foundation of all our work in the OHDSI community, and I'm proud that our open community data standard has been so widely adopted and so extensively used to generate reliable evidence."

- Clair Blacketer
2020 Titan Award for Data Standards recipient



OMOP CDM By The Numbers

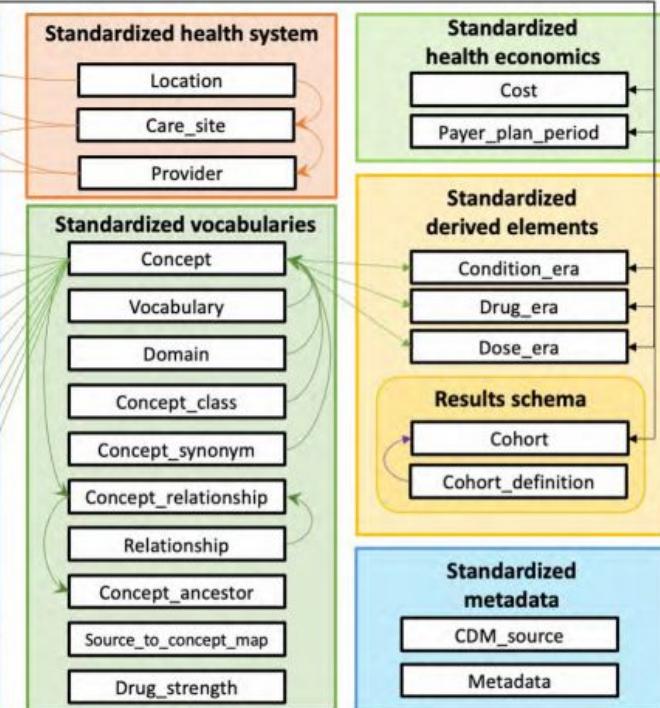
37 tables

- 17 to standardize clinical data
- 10 to standardize vocabularies

394 fields

- 193 with `id` to standardize identification
- 101 with `concept_id` to standardize content
- 43 with `source_value` to preserve original data

1 Open Community Data Standard

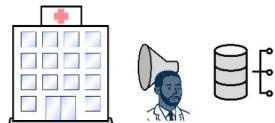


Classical Observational Informatics

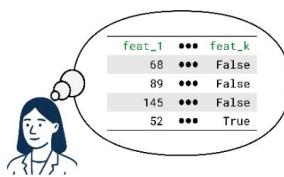
Observational informatics communities draw much more from traditional statistics and informatics backgrounds, and often have much greater clinical expertise and more targeted health problems in mind.



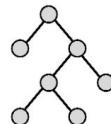
Cohorts in observational informatics studies are typically narrower, and designed via expert input to target the specific problem of interest.



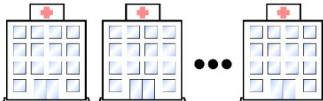
Features are often selected via manual feature engineering efforts to maximize performance, explainability, and robustness for low capacity models.



Models are often lower capacity in traditional observational informatics studies, providing generalizability, robustness, and alignment to understandable clinical mechanisms at the expense of performance.

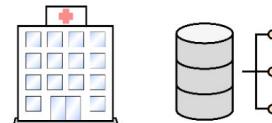
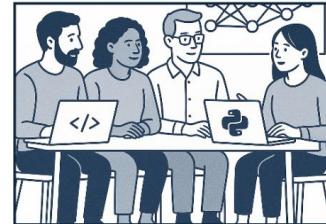


Specific analytic findings or models are expected to generalize across sites of care by virtue of causal, biological grounding.

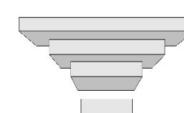
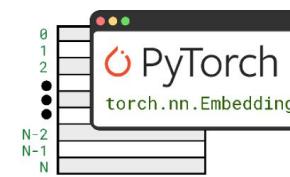


Health AI

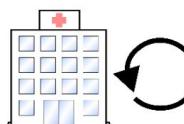
Health AI communities draw heavily from the computer science & AI communities, and leverage the well developed Python ecosystem for AI development in their work. Their research often focuses more on method development, without a concrete clinical target in mind.



Health AI studies often leverage as much data in as close to the raw form as is possible, so the model can learn relationships directly from the source, across all patients.



Health AI models typically omit manual feature engineering, and represent features via learned embedding layers that are randomly initialized, updated through model training, and cover almost all features in the data.



Models are generally high capacity, with model capacity often only limited by data and computational resource availability. Higher capacity models offer higher performance, but are less explainable and more specialized to local data context.

As models are locally specialized and may be too sensitive for release, models are often best for local use, though as centralized data availability increases, this may change.

MEDS-EIC-AR Model

[main](#)[10 Branches](#)[13 Tags](#)[Go to file](#)[Add file](#)[Code](#)

 mmcdermott	Merge pull request #63 from mmcdermott/codex/fix-issues-...		c36f1dd · 2 weeks ago	148 Commits
 .github/workflows	Updated PR action. Closes #44.		2 weeks ago	
 src/MEDS_EIC_AR	Added datamodule config mod for generation to genera...		2 weeks ago	
 tests	Add non-empty check for generated trajectory DataFram...		2 weeks ago	
 .gitignore	Initial commit		2 months ago	
 .pre-commit-config.yaml	Updated pre-commit config to remove shfmt which caus...		2 months ago	
 LICENSE	Initial commit		2 months ago	
 README.md	Fixed tests and removed unused sample ID		2 weeks ago	
 conftest.py	Something's off with dependencies, trying github CI		2 weeks ago	
 pyproject.toml	Something's off with dependencies, trying github CI		2 weeks ago	

[README](#)[MIT license](#)

About

A MEDS, "Everything-is-code" style Autoregressive Generative Model, capable of zero-shot inference.

[Readme](#)[MIT license](#)[Activity](#)[1 star](#)[1 watching](#)[0 forks](#)

Releases 13

[0.1.10](#) (Latest)

2 weeks ago

[+ 12 releases](#)

Packages

No packages published

[Publish your first package](#)

Deployments 13

 pypi 2 weeks ago[+ 12 deployments](#)

Languages

 Python 100.0%

MEDS "Everything-is-code" Autoregressive Model

 v0.1.10  3.12  95%  passing  Main passing  MIT  welcome

contributors 1

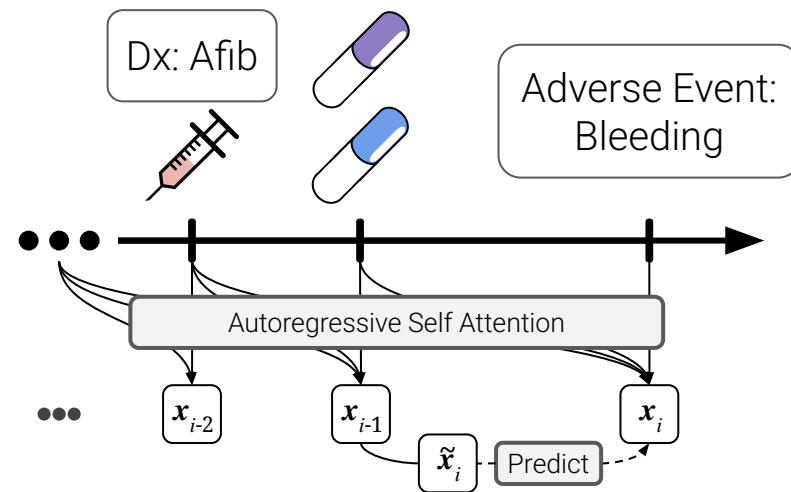
A MEDS, "Everything-is-code" style Autoregressive Generative Model, capable of zero-shot inference.

This is based on the [MEDS-Torch](#) model of the same name.

Installation

```
pip install MEDS-EIC-AR
```

Autoregressive models predict the next MEDS measurement given the patient's history:



Installation and Usage

```
# 1. Install
pip install MEDS-EIC-AR

# 2. Data Pre-processing
MEICAR_process_data input_dir="$RAW_MEDS_DIR" \
    intermediate_dir="$INTERMEDIATE_DIR" \
    output_dir="$FINAL_DATA_DIR"

# 3. Pre-train the model
MEICAR_pretrain datamodule.config.tensorized_cohort_dir="$FINAL_DATA_DIR" \
    output_dir="$PRETRAINED_MODEL_DIR" \
    datamodule.batch_size=32

# 4. Generate Trajectories
MEICAR_generate_trajectories \
    output_dir="$GENERATED_TRAJECTORIES_DIR" \
    model_INITIALIZATION_dir="$PRETRAINED_MODEL_DIR" \
    datamodule.config.tensorized_cohort_dir="$FINAL_DATA_DIR" \
    datamodule.config.task_labels_dir="$TASK_ROOT_DIR/$TASK_NAME" \
    datamodule.batch_size=32
```

The Data Preprocessing

```
# 2. Data Pre-processing
MEICAR_process_data input_dir="$RAW_MEDS_DIR" \
intermediate_dir="$INTERMEDIATE_DIR" \
output_dir="$FINAL_DATA_DIR"
```

```
stages:
  - add_time_derived_measurements:
      timeline_tokens:
        time_unit: years
  - count_codes:
      aggregations:
        - code/n_occurrences
        - code/n_subjects
  - filter_measurements:
      min_subjects_per_code: 1000
  - filter_subjects:
      min_events_per_subject: 10
  - fit_quantile_binning:
      aggregations:
        - name: values/quantiles
          quantiles: [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
  - bin_numeric_values:
      custom_bins:
        TIMELINE//DELTA//years:
          1_minute: 0.00000190258
          5_minutes: 0.00000951293
          10_minutes: 0.00001902587
          30_minutes: 0.00005707762
          ...
          1_year: 1
          2_years: 2
          5_years: 5
          10_years: 10
          20_years: 20
          40_years: 40
```

[README](#) [MIT license](#)

MEDS Transforms: Build and run complex pipelines over MEDS datasets via simple parts

[pypi v0.5.1](#) [Python 3.12](#) [MEDS 0.4](#) [docs passing](#) [codecov 100%](#) [Tests passing](#) [Code Quality Main passing](#)
[Config](#) [Hydra 1.3](#) [License MIT](#) [PRs welcome](#) [contributors 7](#)

🚀 Quick Start

1. Install via pip:

```
pip install MEDS-transforms
```

2. Craft a pipeline YAML file:

```
input_dir: $MEDS_ROOT
output_dir: $PIPELINE_OUTPUT

description: Your special pipeline

stages:
  - filter_subjects:
      min_events_per_subject: 5
  - add_time_derived_measurements:
      age:
        DOB_code: MEDS_BIRTH
        age_code: AGE
        age_unit: years
      time_of_day:
        time_of_day_code: TIME_OF_DAY
        endpoints: [6, 12, 18, 24]
  - fit_outlier_detection:
      _base_stage: aggregate_code_metadata
      aggregations:
        - values/n_occurrences
        - values/sum
        - values/sum_sqd
  - occlude_outliers:
      stddev_cutoff: 1
  - fit_normalization:
      _base_stage: aggregate_code_metadata
      aggregations:
        - code/n_occurrences
        - code/n_subjects
        - values/n_occurrences
        - values/sum
        - values/sum_sqd
  - fit_vocabulary_indices
  - normalization
```

Save your pipeline YAML file on disk at `$PIPELINE_YAML`.

3. Run the pipeline

In the terminal, run

```
MEDS_transform-pipeline pipeline_fp="$PIPELINE_YAML"
```

After you do, you will see output files stored in `$PIPELINE_OUTPUT` with the results of each stage of the pipeline, stored in stage specific directories, and the global output in `$PIPELINE_OUTPUT/data` and `$PIPELINE_OUTPUT/metadata` (for data and metadata outputs, respectively). That's it!

The Model

```
# 3. Pre-train the model
MEICAR_pretrain datamodule.config.tensorized_cohort_dir="
    output_dir="$PRETRAINED_MODEL_DIR" \
    datamodule.batch_size=32
```

```
import torch
import torch.nn.functional as F
from meds_torchdata import MEDSTorchBatch
from omegaconf import DictConfig
from transformers import (
    AutoConfig,
    AutoModelForCausalLM,
    GenerationConfig,
    GPTNeoXConfig,
    GPTNeoXForCausalLM,
)
from transformers.modeling_outputs import CausalLMOutputWithPast

class Model(torch.nn.Module):
    def __init__(self, gpt_kwargs: dict | DictConfig, precision: str = "32-true", do_demo: bool = False):
        super().__init__()

        self.HF_model_config: GPTNeoXConfig = AutoConfig.from_pretrained("EleutherAI/gpt-neox-20b")
        ...
        self.HF_model = AutoModelForCausalLM.from_config(self.HF_model_config, **extra_kwargs)

    def _hf_inputs(self, batch: MEDSTorchBatch) -> dict[str, torch.Tensor]:
        return {"input_ids": batch.code, "attention_mask": (batch.code != batch.PAD_INDEX)}

    def forward(self, batch: MEDSTorchBatch) -> tuple[torch.FloatTensor, CausalLMOutputWithPast]:
        outputs = self.HF_model(**self._hf_inputs(batch))
        loss = F.cross_entropy(
            outputs.logits[:, :-1].transpose(2, 1), batch.code[:, 1:], ignore_index=batch.PAD_INDEX
        )
        return loss, outputs

    def generate(self, batch: MEDSTorchBatch, do_sample: bool = True, **kwargs) -> torch.Tensor:
        for_hf = self._hf_inputs(batch)

        generation_config = GenerationConfig(
            max_new_tokens=self.max_seq_len - batch.code.shape[1],
            do_sample=do_sample,
            num_beams=1, # no beam search
            temperature=1.0,
            pad_token_id=batch.PAD_INDEX,
            bos_token_id=None,
            eos_token_id=self.HF_model.config.eos_token_id,
        )

        output_ids = self.HF_model.generate(
            for_hf.pop("input_ids"),
            generation_config=generation_config,
            **for_hf,
            **kwargs,
        )
        input_seq_len = batch.code.shape[1]
        return output_ids[:, input_seq_len:]
```

The Model

```
# 3. Pre-train the model
MEICAR_pretrain datamodule.config.tensorized_cohort_dir="
    output_dir="$PRETRAINED_MODEL_DIR" \
    datamodule.batch_size=32
```

```
import torch
import torch.nn.functional as F
from medstorchdata import MEDSTorchBatch
from omegaconf import DictConfig
from transformers import (
    AutoConfig,
    AutoModelForCausalLM,
    GenerationConfig,
    GPTNeoXConfig,
    GPTNeoXForCausalLM,
)
from transformers.modeling_outputs import CausalLMOutputWithPast

class Model(torch.nn.Module):
    def __init__(self, gpt_kwarg: dict | DictConfig, precision: str = "32-true", do_demo: bool = False):
        super().__init__()

        self.HF_model_config: GPTNeoXConfig = AutoConfig.from_pretrained("EleutherAI/gpt-neox-20b")
        ...
        self.HF_model = AutoModelForCausalLM.from_config(self.HF_model_config, **extra_kwarg)

    def _hf_inputs(self, batch: MEDSTorchBatch) -> dict[str, torch.Tensor]:
        return {key: value for key, value in batch.items() if key != "attention_mask": (batch.code != batch.PAD_INDEX)}

    def forward(self, batch: MEDSTorchBatch) -> tuple[torch.FloatTensor, CausalLMOutputWithPast]:
        outputs = self.HF_model(**self._hf_inputs(batch))
        loss = F.cross_entropy(
            outputs.logits[:, :-1].transpose(2, 1), batch.code[:, 1:], ignore_index=batch.PAD_INDEX
        )
        return loss, outputs

    def generate(self, batch: MEDSTorchBatch, do_sample: bool = True, **kwargs) -> torch.Tensor:
        for_hf = self._hf_inputs(batch)

        generation_config = GenerationConfig(
            max_new_tokens=self.max_seq_len - batch.code.shape[1],
            do_sample=do_sample,
            num_beams=1, # no beam search
            temperature=1.0,
            pad_token_id=batch.PAD_INDEX,
            bos_token_id=None,
            eos_token_id=self.HF_model.config.eos_token_id,
        )

        output_ids = self.HF_model.generate(
            for_hf.pop("input_ids"),
            generation_config=generation_config,
            **for_hf,
            **kwargs,
        )
        input_seq_len = batch.code.shape[1]
        return output_ids[:, input_seq_len:]
```

The

```
# 3. Pre-train  
MEICAR_pretrain  
output_dir=  
datamodule.
```

```
MEDSTorchBatch:  
| Mode: Subject-Measurement (SM)  
| Static data? ✓  
| Labels? ✗  
  
Shape:  
| Batch size: 2  
| Sequence length: 5  
  
All dynamic data: (2, 5)  
Static data: (2, 2)  
  
Data:  
| Dynamic:  
| | time_delta_days (torch.float32):  
| | | [[0.00e+00, 1.18e+04, ..., 0.00e+00, 9.79e-02],  
| | | [0.00e+00, 1.24e+04, ..., 0.00e+00, 4.64e-02]]  
| | code (torch.int64):  
| | | [[ 5,  3, ..., 11,  4],  
| | | [ 5,  2, ..., 11,  4]]  
| | numeric_value (torch.float32):  
| | | [[ 0.00,  0.00, ..., -0.34,  0.00],  
| | | [ 0.00,  0.00, ...,  0.85,  0.00]]  
| | numeric_value_mask (torch.bool):  
| | | [[False, False, ..., True, False],  
| | | [False, False, ..., True, False]]  
  
| Static:  
| | static_code (torch.int64):  
| | | [[8, 9],  
| | | [8, 9]]  
| | static_numeric_value (torch.float32):  
| | | [[ 0.00, -0.54],  
| | | [ 0.00, -1.10]]  
| | static_numeric_value_mask (torch.bool):  
| | | [[False,  True],  
| | | [False,  True]]
```

```
import torch  
import torch.nn.functional as F  
from medstorchdata import MEDSTorchBatch  
from omegaconf import DictConfig  
from transformers import (  
    AutoConfig,  
    AutoModelForCausalLM,  
    GenerationConfig,  
    GPTNeoXConfig,  
    GPTNeoXForCausalLM,  
)  
from transformers.modeling_outputs import CausalLMOutputWithPast  
  
class Model(torch.nn.Module):  
    def __init__(self, **kwargs: dict | DictConfig, precision: str = "32-true", do_demo: bool = False):  
        super().__init__()  
  
        self.HF_model_config: GPTNeoXConfig = AutoConfig.from_pretrained("EleutherAI/gpt-neo-20b")  
  
        ...  
  
        self.HF_model = AutoModelForCausalLM.from_config(self.HF_model_config, **extra_kwargs)  
  
    def _hf_inouts(self, batch: MEDSTorchBatch) -> dict[str, torch.Tensor]:  
        return {  
            "input_ids": batch.code, "attention_mask": (batch.code != batch.PAD_INDEX)}  
  
    def forward(self, batch: MEDSTorchBatch) -> tuple[torch.FloatTensor, CausalLMOutputWithPast]:  
        outputs = self.HF_model(**self._hf_inouts(batch))  
        loss = F.cross_entropy(outputs.logits[:, :-1].transpose(0, 1), batch.code[:, 1:], ignore_index=batch.PAD_INDEX)  
  
        return loss, outputs  
  
    def generate(self, batch: MEDSTorchBatch, do_sample: bool = True, **kwargs) -> torch.Tensor:  
        for_hf = self._hf_inouts(batch)  
  
        generation_config = GenerationConfig(  
            max_new_tokens=self.max_seq_len - batch.code.shape[1],  
            do_sample=do_sample,  
            num_beams=1, # no beam search  
            temperature=1.0,  
            pad_token_id=batch.PAD_INDEX,  
            bos_token_id=None,  
            eos_token_id=self.HF_model.config.eos_token_id,  
        )  
  
        output_ids = self.HF_model.generate(  
            for_hf.pop("input_ids"),  
            generation_config=generation_config,  
            **for_hf,  
            **kwargs,  
        )  
  
        input_seq_len = batch.code.shape[1]  
  
        return output_ids[:, :input_seq_len:]
```

MEDS TorchData: A PyTorch Dataset Class for MEDS Datasets

 PyTorch  Python 3.11+  pypi v0.6.1  MEDS 0.4  docs passing  Tests passing  codecov 100%

 Code Quality Main passing  Config Hydra 1.3  contributors 7  PRs welcome  License MIT

🚀 Quick Start

Step 1: Install

```
pip install meds-torch-data
```



Step 2: Data Tensorization

⚠ Warning

If your dataset is not sharded by split, you need to run a reshard to split stage first! You can enable this by adding the `do_reshard=True` argument to the command below.

If your input MEDS dataset lives in `$MEDS_ROOT` and you want to store your pre-processed files in `$PYD_ROOT`, you run:

```
MTD_preprocess MEDS_dataset_dir="$MEDS_ROOT" output_dir="$PYD_ROOT"
```



Step 3: Use the dataset

To use a dataset, you need to (1) define your configuration object and (2) create the dataset object. The only required configuration parameters are `tensorized_cohort_dir`, which points to the root directory containing the pre-processed data on disk (`$PYD_ROOT` in the above example), and `max_seq_len`, which is the maximum sequence length you want to use for your model. Here's an example:

```
import os
from meds_torchdata import MEDSPytorchDataset, MEDSTorchDataConfig

cfg = MEDSTorchDataConfig(tensorized_cohort_dir=os.environ["PYD_ROOT"], max_seq_len=512)
pyd = MEDSPytorchDataset(cfg, split="train")
```



stages:

- `fit_normalization`:
- `aggregations`:
 - "code/n_occurrences"
 - "code/n_subjects"
 - "values/n_occurrences"
 - "values/sum"
 - "values/sum_sqd"
- `fit_vocabulary_indices`
- `normalization`
- `tokenization`
- `tensorization`

If your input MEDS dataset lives in `$MEDS_ROOT` and you want to store your pre-processed files in `$PYD_ROOT`, you run:

```
MTD_preprocess MEDS_dataset_dir="$MEDS_ROOT" output_dir="$PYD_ROOT"
```

Step 3: Use the dataset

To use a dataset, you need to (1) define your configuration object and (2) create the dataset object. The only required configuration parameters are `tensorized_cohort_dir`, which points to the root directory containing the pre-processed data on disk (`$PYD_ROOT` in the above example), and `max_seq_len`, which is the maximum sequence length you want to use for your model. Here's an example:

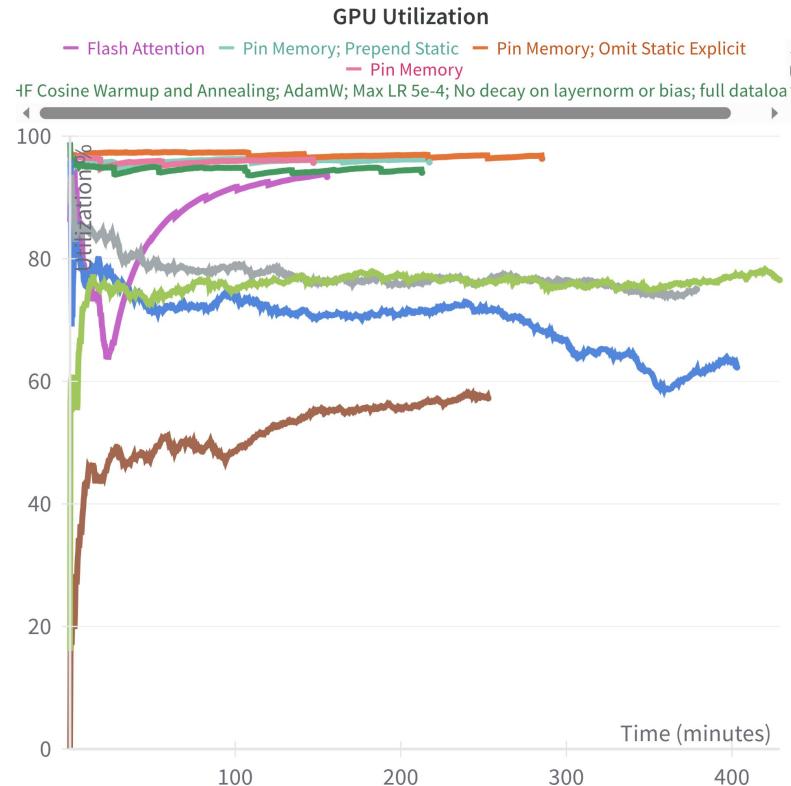
```
import os
from meds_torchdata import MEDSPytorchDataset, MEDSTorchDataConfig

cfg = MEDSTorchDataConfig(tensorized_cohort_dir=os.environ["PYD_ROOT"], max_seq_len=512)
pyd = MEDSPytorchDataset(cfg, split="train")
```

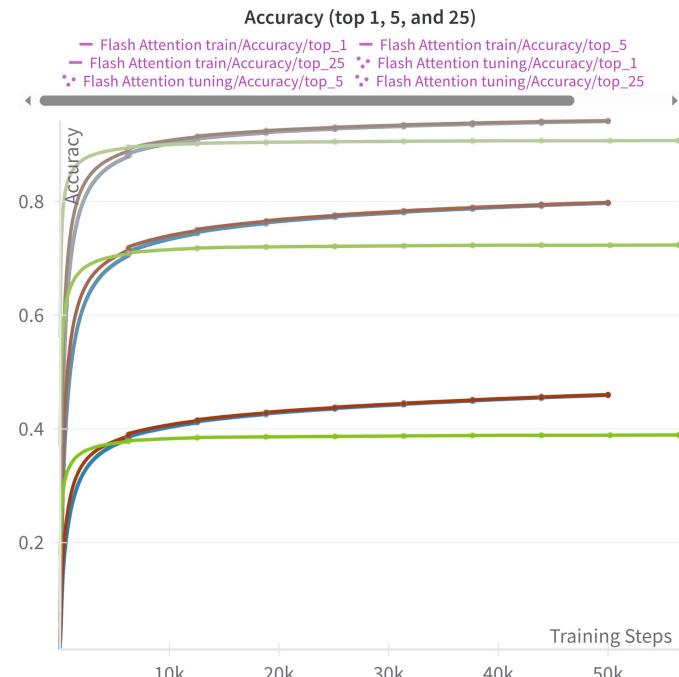
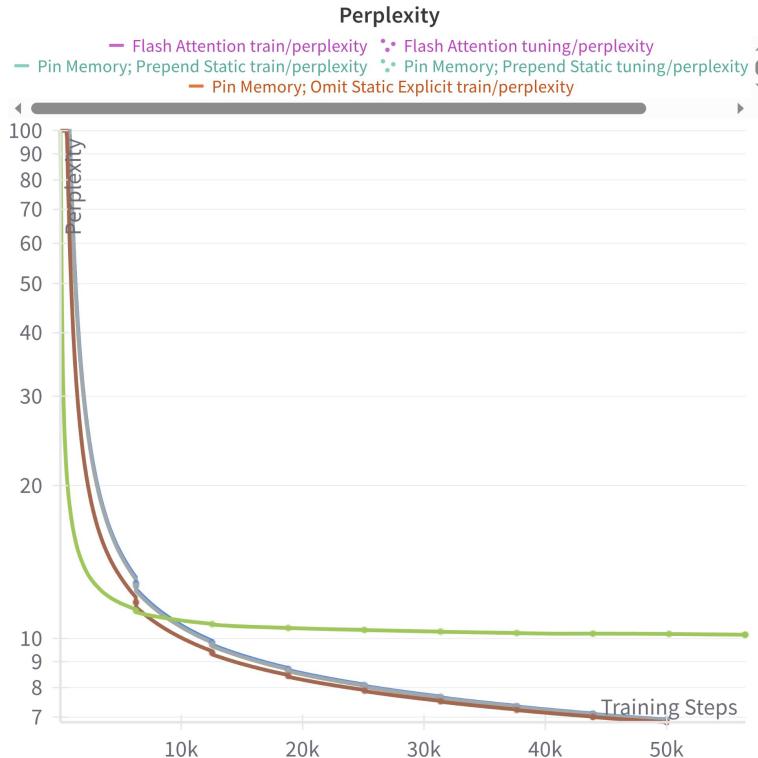
Batches

```
● ○ ●  
MEDSTorchBatch:  
| Mode: Subject-Measurement (SM)  
| Static data? ✓  
| Labels? ✗  
  
Shape:  
| Batch size: 2  
| Sequence length: 5  
  
All dynamic data: (2, 5)  
Static data: (2, 2)  
  
Data:  
| Dynamic:  
| | time_delta_days (torch.float32):  
| | | [[0.00e+00, 1.18e+04, ..., 0.00e+00, 9.79e-02],  
| | | | [0.00e+00, 1.24e+04, ..., 0.00e+00, 4.64e-02]]  
| | code (torch.int64):  
| | | [[5, 3, ..., 11, 4],  
| | | | [5, 2, ..., 11, 4]]  
| | numeric_value (torch.float32):  
| | | [[0.00, 0.00, ..., -0.34, 0.00],  
| | | | [0.00, 0.00, ..., 0.85, 0.00]]  
| | numeric_value_mask (torch.bool):  
| | | [[False, False, ..., True, False],  
| | | | [False, False, ..., True, False]]  
  
Static:  
| static_code (torch.int64):  
| | [[8, 9],  
| | | [8, 9]]  
| static_numeric_value (torch.float32):  
| | [[0.00, -0.54],  
| | | [0.00, -1.10]]  
| static_numeric_value_mask (torch.bool):  
| | [[False, True],  
| | | [False, True]]
```

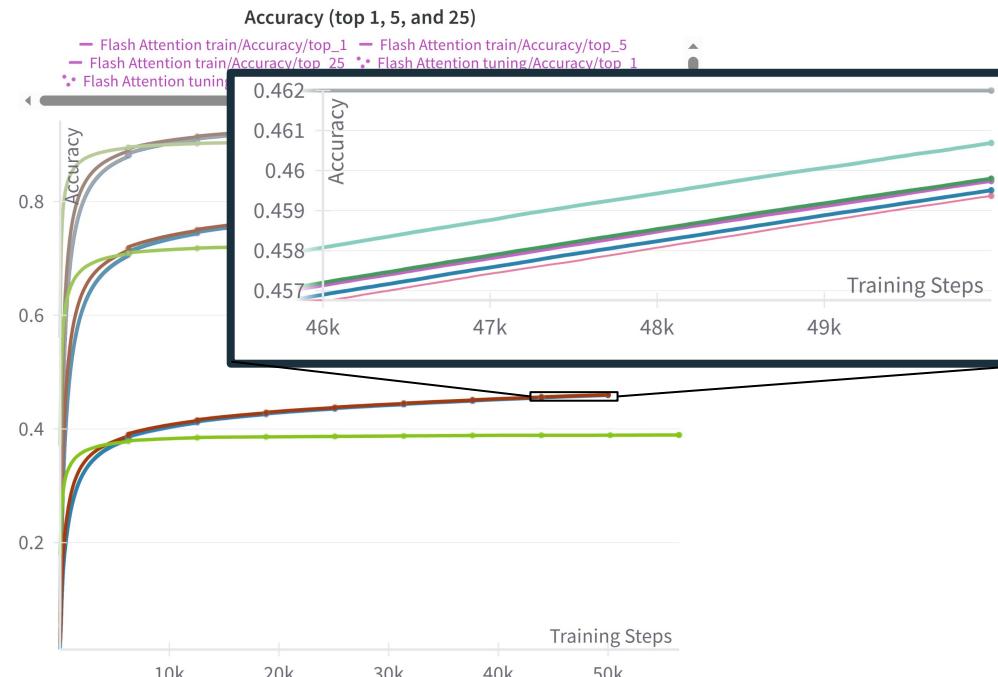
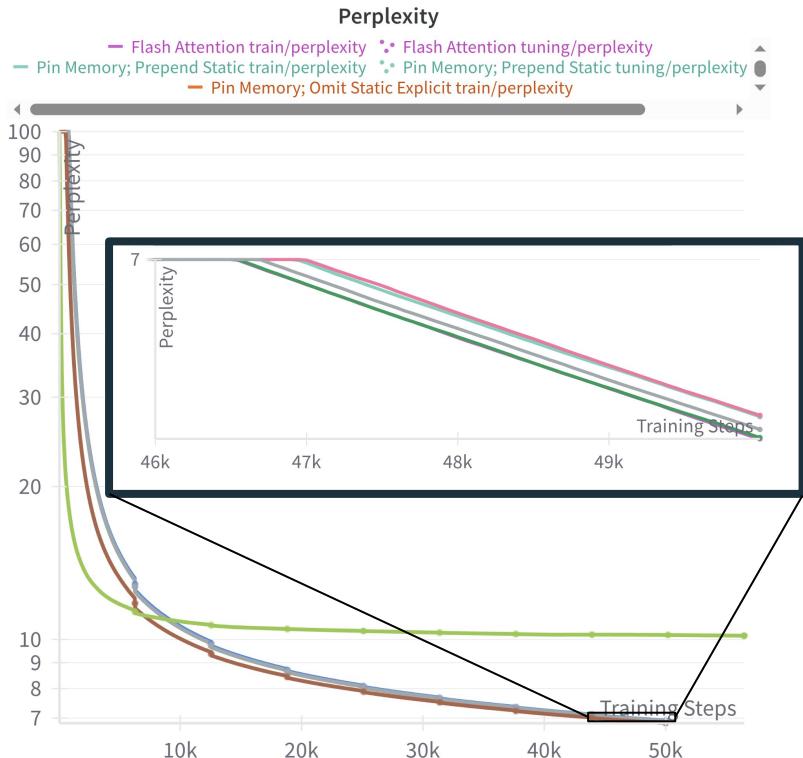
GPU Utilization (Training Speed)



Training Convergence and Stability



Training Convergence and Stability



Determining Parameters?



Determining Parameters?

```
>>> print_directory("./src/MEDS_EIC_AR/configs", config=PrintConfig(file_extension=".yaml"))
├── _demo_generate_trajectories.yaml
├── _demo_pretrain.yaml
└── _generate_trajectories.yaml
    └── _pretrain.yaml
        ├── datamodule
        │   ├── default.yaml
        │   └── generate_trajectories.yaml
        └── pretrain.yaml
            ├── inference
            │   ├── default.yaml
            │   └── demo.yaml
            └── lightning_module
                ├── LR_scheduler
                │   ├── cosine_annealing_warm_restarts.yaml
                │   ├── get_cosine_schedule_with_warmup.yaml
                │   ├── one_cycle_LR.yaml
                │   └── reduce_LR_on_plateau.yaml
                ├── default.yaml
                └── demo.yaml
                └── metrics
                    └── default.yaml
                └── model
                    ├── default.yaml
                    ├── demo.yaml
                    └── small.yaml
                └── optimizer
                    ├── adam.yaml
                    └── adamw.yaml
            └── trainer
                ├── callbacks
                │   ├── default.yaml
                │   ├── early_stopping.yaml
                │   ├── learning_rate_monitor.yaml
                │   └── model_checkpoint.yaml
                ├── default.yaml
                └── demo.yaml
                └── logger
                    ├── csv.yaml
                    ├── mlflow.yaml
                    └── wandb.yaml
```

Other Tool Posters



MEDS-Reader

A blazing fast EHR processing library

Ethan Steinberg, Michael Worowin, Suhana Bedi, Jason Alan Fries, Matthew B. A. McDermott, Nigam Shah



Introduction

With the rise of foundation models, EHR research needs to process increasing amounts of data. However, existing tooling, which is optimized for tabular data, images, and text, has not kept up.

To fix this we introduce `meds_reader`, a blazing fast EHR processing library.

Contributions

We introduce `meds_reader`, which has:

- A simple Python interface, easily installable with pip
- A flexible interface, usable for
 - Featurization, labeling, tensorization, and much more!
- 10-100x speed and memory improvements over alternatives
- Support for a variety of EHR datasets through the MEDS standard

EHRs as Event Streams

EHR Dataset = Collection of subjects

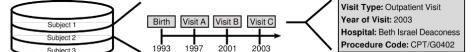
EHR Dataset

Subject = Series of events

Subject

Event = Collection of properties

Event



Python API

Core Design:
Subject transformation functions

Features:

- Parallelization across threads
- Uses incremental algorithms
- Has a C++ backend
- Supports arbitrary properties
- Memory and disk compression

```
def find_stroke(subjects):
    subject_ids = set()
    for subject in subjects:
        for event in subject.events:
            if event.code.startswith('ICD10/163'):
                subject_ids.add(subject.subject_id)
    return subject_ids

subject_database = meds_reader.SubjectDatabase('...', num_threads=10)
all_subjects = set()
for partial_result in subject_database.map(find_stroke):
    all_subjects |= partial_result
```

`pip install meds_reader`

Benchmarks

We can test the performance of `meds_reader` by reimplementing two other data processing pipelines.

Dataset	#Subjects	Library	Memory (GB)	CPU (sec)	Disk (GB)
eICU	200,859	PyHealth	16.80	0.67	23.1
MIMIC III	46,720	PyHealth	20.97	0.54	38.8
MIMIC-IV	299,712	MDP	11.32	2.08	6.0

Table 1: Compute resources used by various data processing pipelines compared to our reimplementations using `meds_reader`. We measure peak memory usage, wall clock CPU time, and the disk space used for each direction. We compare the original resources used to the reimplementations and report the improvement factor (A-B). Orig = Original pipeline, MDP = MIMIC-IV Data Pipeline.

We observe huge improvements in memory, speed and disk usage.

Dataset Support

Supports any dataset has been converted to MEDS (Medical Event Data Standard)

OMOP

MIMIC-IV

eICU

EHRSHOT

Your Custom Dataset

[README](#) [MIT license](#)

MEDS Transforms: Build and run complex pipelines over MEDS datasets via simple parts

[pip v0.5.1](#) [Python 3.12](#) [MEDS 0.4](#) [docs](#) [passing](#) [codecov](#) [100%](#) [Tests](#) [passing](#) [Code Quality Main](#) [passing](#)
[Config](#) [Hydra 1.3](#) [License](#) [MIT](#) [PRs](#) [welcome](#) [contributors](#) 7

Quick Start

1. Install via pip:

`pip install MEDS-transforms`

2. Craft a pipeline YAML file:

```
input_dir: $MEDS_ROOT
output_dir: $PIPELINE_OUTPUT
```

```
description: Your special pipeline
```

stages:

- filter_subjects:
 - min_events_per_subject: 5
- add_time_derived_measurements:
 - age:
 - DOB_code: MEDS_BIRTH
 - age_code: AGE
 - age_unit: years
 - time_of_day:
 - time_of_day_code: TIME_OF_DAY
 - endpoints: [6, 12, 18, 24]
- fit_outlier_detection:
 - base_stage: aggregate_code_metadata
 - aggregations:
 - values/n_occurrences
 - values/sum
 - values/sum_sqd
 - occlude_outliers:
 - stddev_cutoff: 1
- fit_normalization:
 - base_stage: aggregate_code_metadata
 - aggregations:
 - code/n_occurrences
 - code/n_subjects
 - values/n_occurrences
 - values/sum
 - values/sum_sqd
 - fit_vocabulary_indices
 - normalization

MEDS-Torch: A Machine Learning Pipeline for Inductive Experiments on EHR Data

Nassim Oufatolle, Teya Bergamaschi, Paweł Renc, Alekcia Kolo, Matthew B.A. McDermott, Collin Stultz

Streamlining Health AI Benchmarking

Challenges of Reproducible benchmarking across health AI systems

1. Complex data preprocessing requirements (\uparrow User Burden)
2. Efficient loading of large irregularly sampled Time Series data

MEDS-Torch Key Features

- Multiple tokenization and pretraining methods to enable benchmarking
- Generality across datasets via the MEDS data standard
- Minimal code and user burden
- GPU-optimized processing via MEDS-torch Tensorization
- Memory and CPU efficient Dataloaders for contrastive learning, supervised, and forecasting workflows

MEDS-Torch Usage

Stage 1 (Tensorization) What your need:
 • Data in MEDS format
 ➤ Task labels

Stage 2 (Training/Prediction) What you need:
 • Tokenized Data
 ➤ Task labels

MEDS-Torch Results

Tokenization Performance Comparison

Method Performance Comparison

The Science of ML in Healthcare

Health AI faces a systemic reproducibility crisis, limiting our ability to do effective science. The MEDS health AI ecosystem changes that.

- MEDS enables a robust ecosystem of health AI tools.
- This allows groups to share software, models, & results!
- Check out the link below to learn more!

To Learn More About MEDS Visit: [https://meds.ai](#)

Future Roadmap

- Multimodal Support (+ ECGs, Notes, etc.)
- Zero-Shot Autoregressive Generation and Evaluations
- CPU and Memory Benchmarking Scripts

[README](#) [License](#)

MEDS TorchData: A PyTorch Dataset Class for MEDS Datasets

PyTorch Python 3.11+ pypi v0.6.1 MEDS 0.4 docs passing Tests passing codecov 100%
 Code Quality Main passing Config Hydra 1.3 contributors 7 PRs welcome License MIT

Quick Start

Step 1: Install

```
pip install medstorch-data
```

Step 2: Data Tensorization

⚠ Warning

If your dataset is not sharded by split, you need to run a reshard to split stage first! You can enable this by adding the `do_reshard=True` argument to the command below.

If your input MEDS dataset lives in `$MEDS_ROOT` and you want to store your pre-processed files in `$PYD_ROOT`, you run:

```
MTD_preprocess MEDS_dataset_dir="$MEDS_ROOT" output_dir="$PYD_ROOT"
```

Step 3: Use the dataset

To use a dataset, you need to (1) define your configuration object and (2) create the dataset object. The only required configuration parameters are `tensorized_cohort_dir`, which points to the root directory containing the pre-processed data on disk (`$PYD_ROOT` in the above example), and `max_seq_len`, which is the maximum sequence length you want to use for your model. Here's an example:

```
import os
from medstorchdata import MEDSPytorchDataset, MEDSTorchDataConfig

cfg = MEDSTorchDataConfig(tensorized_cohort_dir=os.environ["PYD_ROOT"], max_seq_len=512)
pyd = MEDSPytorchDataset(cfg, split="train")
```

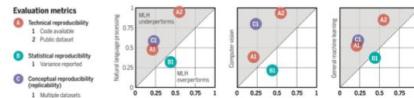
MEDS-DEV: Enabling Reproducible Science of AI for EHR Data

Aleksia Kolo, Chao Pang, Edward Choi, Ethan Steinberg, Hyewon Jeong, Jack Gallifant, Jason A. Fries, Jeffrey N. Chiang, Jungwoo Oh, Justin Xu, Kamilé Stankevičiūtė, Kiril V. Klein, Matthew McDermott, Mikkel Odgaard, Nassim Oufatolle, Patrick Rockenschaub, Paweł Renc, Robin van de Water, Shalmali Joshi, Simon A. Lee, Teya S. Bergamaschi, Tom J. Pollard, Vincent Jeanseime, Young Sang Choi



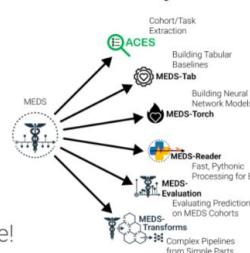
Health AI has a Reproducibility Crisis

Health AI faces a systemic reproducibility crisis, limiting our ability to do effective science. The MEDS health AI ecosystem, including ACES and MEDS-DEV, changes that.

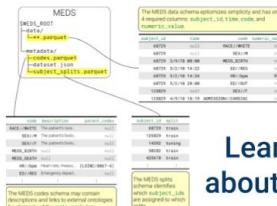


MEDS: A Health AI Data Ecosystem

- MEDS enables a robust ecosystem of health AI tools.
- This allows groups to share software, models, & results!
- Check out the link below to learn more!



- MEDS is simple and easy to use.
- Converting to MEDS is easy!



Learn More about MEDS:

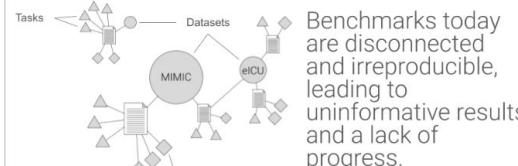


MEDS-DEV Decentralized Extensible Validation

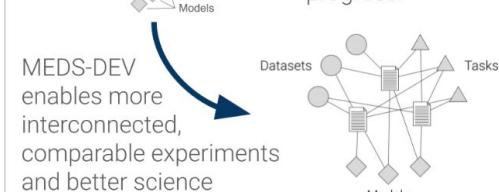
Meaningful shared tasks in healthcare require:

- Decentralized evaluation over private data.
- *Frictionlessly reproducible* task definitions, data preprocessing pipelines, model training recipes, and evaluation protocols.
- A community curated set of tasks spanning clinical problems, settings, and data needs.

MEDS-DEV enables meaningful benchmarking and reproducible AI in ML4H by satisfying these needs.



Benchmarks today are disconnected and irreproducible, leading to uninformative results and a lack of progress.



How do you use MEDS-DEV?

Need	Tool
Reproducible, transparent task definitions	ACES task configs are transparent, reproducible, & can be shared across datasets
Shared evaluation protocols	MEDS-Evaluation ensures evaluation is consistent and aspects such as fairness and equity can be included
Communal contributions on tasks & models	MEDS-DEV configuration files and results are stored & versioned in an open-source GitHub, enabling communal development and decentralization of results
Decentralized, extensible data and task support	MEDS-DEV
	Model MIMIC-IV/Columbia Long LOS Model ICU Mortality
	Log. Reg. 0.752/0.677 0.754/0.509 Light GBM 0.783/0.757 0.798/0.661 MOTOR 0.804/0.735 0.854/0.727 CEHR-BERT 0.808/0.741 0.845/0.726 MEDS-Tab 0.811 / 0.761 0.830 / 0.785 GenHPF 0.779/0.662 0.790/0.633

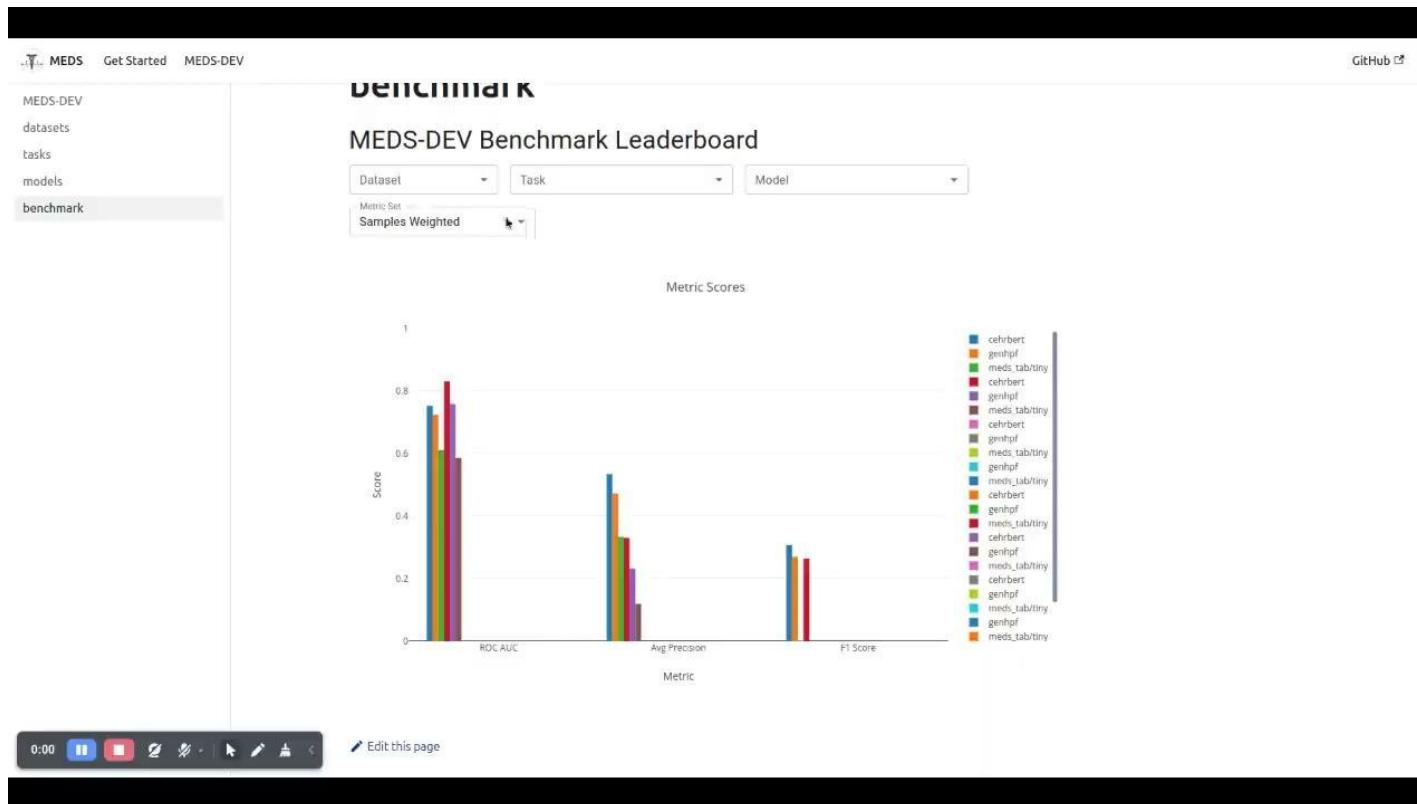
Check out detailed MEDS-DEV Tutorials Here:



Acknowledgements

MEDS gratefully acknowledges support from a Berkowitz Postdoctoral Fellowship. KS acknowledges support from Astrazeneca. JAF acknowledges support for this effort from the Debra and Mark Leslie fund for supporting AI in Healthcare research. T2 is supported by the National Institutes of Health (NIH) 01702632701 and NIH R01MH117500. The work of Dr. Chih-Jen Lin was supported by the IITP grant (No. 2019-0700-0190-005) and the NRF grant (NRF-2020R1D1A3A02045465) funded by the Korean Ministry of Science and ICT and the PIA40 project funded by the Innovation Fund Denmark. NHS acknowledges support for this effort from the Debra and Mark Leslie fund for supporting AI in Healthcare research/Pioneer Centre for AI, DNRF Grant P1, Novo Nordisk Foundation

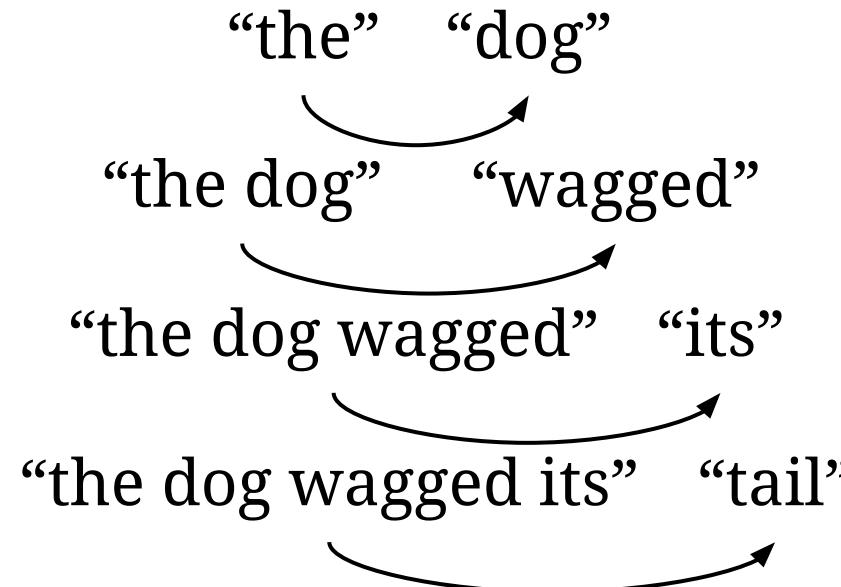
Sharing Results



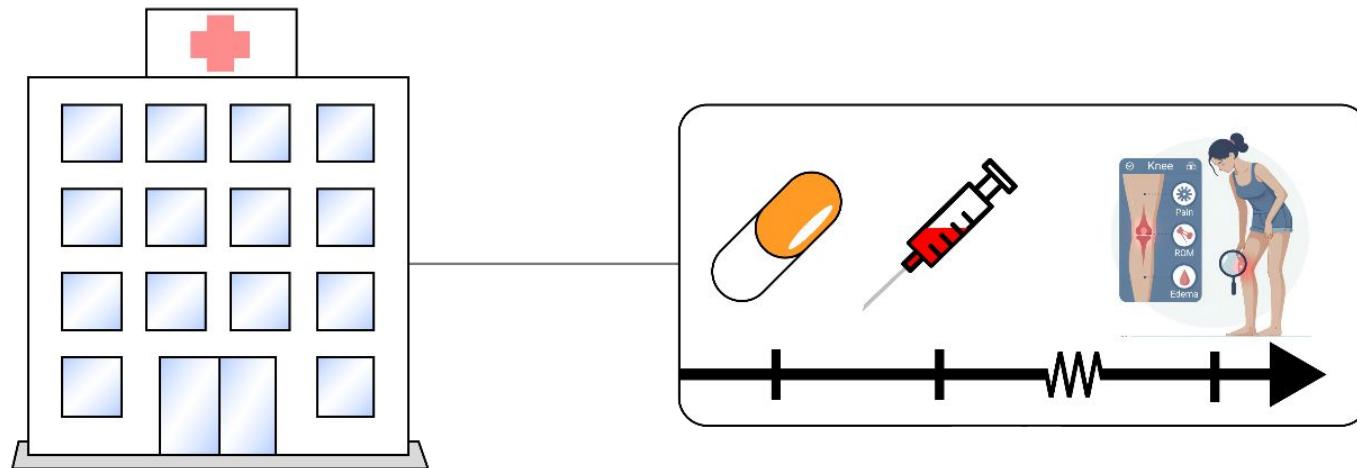
Language Models: What word comes next?

Language Models: What word comes next?

“the dog wagged its tail”
“the”



EHR Data is a Continuous Longitudinal Timeline of Complex Observations



Autoregressive Foundation Models: What observation comes next?

Such models have been widely validated in the literature

CONTEXT CLUES: EVALUATING LONG CONTEXT MODELS FOR CLINICAL PREDICTION TASKS ON EHRs

Michael Wornow^{*1}, Suhana Bedi^{*1}, Miguel Angel Fuentes Hernandez¹, Ethan Steinberg¹², Jason Alan Fries¹, Christopher Re¹, Sanmi Koyejo¹, Nigam H. Shah¹

¹Stanford University ²Prealize Health

ABSTRACT

Foundation Models (FMs) trained on Electronic Health Records (EHRs) have achieved state-of-the-art results on numerous clinical prediction tasks. However, most existing EHR FMs have context windows of <1k tokens. This prevents them from modeling full patient EHRs which can exceed 10k's of events. Recent advancements in subquadratic long-context architectures (e.g., Mamba) of-

Generative Medical Event Models Improve with Scale

Shane Waxler^{*1}, Paul Blazek^{*1}, Davis White^{*1}, Daniel Schneider^{*1}, Kevin Chung¹, Mani Nagarathnam¹, Patrick Williams¹, Hank Voeller¹, Karen Wong¹, Matthew Swanhorst¹, Sheng Zhang², Naoto Usuyama², Cliff Wong², Tristan Naumann², Hoifung Poon², Andrew Loza³, Daniella Meeker^{3, 4}, Seth Hain¹, and Rahul Shah^{†1}

¹Epic Systems

²Microsoft Research

³Yale School of Medicine

⁴Cosmos Governing Council

Abstract

Realizing personalized medicine at scale calls for methods that distill insights from longitudinal patient journeys, which can be viewed as a sequence of medical events. Foundation models pretrained on large-scale medical event data represent a promising direction for scaling real-world evidence generation and generalizing to diverse downstream tasks. Using Epic Cosmos, a dataset with medical events from de-identified longitudinal health records for 16.3 billion encounters over 300 million unique patient records from 310 health systems, we introduce the Comet models, a family of decoder-only transformer models pretrained on 118 million patients representing 115 billion discrete medical events (151 billion tokens). We present the largest scaling-law study of medical event data, establishing a methodology for pretraining and revealing power-law scaling relationships for compute, tokens, and model size. Consequently, we pretrained a series of compute-optimal models with up to 1 billion parameters. Conditioned on a patient's real-world history, Comet autoregressively predicts the next medical event to simulate patient health timelines. We studied 78 real-world tasks, including diagnosis prediction, disease prognosis, and healthcare operations. Remarkably for a foundation model with generic pretraining and simulation-based inference, Comet

npj | digital medicine

Published in partnership with Seoul National University Bundang Hospital

Article



<https://doi.org/10.1038/s41746-024-01235-0>

Zero shot health trajectory prediction using transformer

Check for updates

Pawel Renc ^{1,2,3}, Yugang Jia ⁴, Anthony E. Samir ^{1,2}, Jaroslaw Was ³, Quanzheng Li ^{1,2}, David W. Bates ^{2,3,5} & Arkadiusz Sitek ^{1,2}



Foundation model of electronic medical records for adaptive risk estimation

GigaScience, 2025, 14, 1–12

DOI: 10.1093/gigascience/giaf107

Research

Pawel Renc ^{1,2,3}, Michal K. Grzeszczyk ^{2,3}, Nassim Oufatolle ⁴, Deirdre Goode ^{3,5}, Yugang Jia ⁶, Szymon Bieganski ⁷, Matthew B.

A. McDermott ⁸, Jaroslaw Was ¹, Anthony E. Samir ^{2,3}, Jonathan W. Cunningham ^{3,5}, David W. Bates ^{3,10,11}, and

Arkadiusz Sitek ^{1,2,3}.

¹AGH University of Krakow, Department of Applied Computer Science, al. Mickiewicza 30, 30-059 Kraków, Poland

²Massachusetts General Hospital, Department of Radiology, 55 Fruit St, Suite 427, Boston, MA 02114, USA

³Harvard Medical School, 25 Shattuck Street Boston, MA 02115, USA

⁴Massachusetts Institute of Technology, Electrical Engineering and Computer Science (EECS), 143 Albany St, Cambridge, MA 02139 Unit 133B, USA

⁵Newton Wellesley Hospital, Emergency Department, 2014 Washington St, Newton, MA 02462, USA

⁶Massachusetts Institute of Technology, Laboratory for Computational Physiology, Institute for Medical Engineering and Science, Building E25-505 77 Massachusetts Avenue Boston, MA 02139, USA

⁷Central Clinical Hospital of the Medical University in Lodz, Cardiology - Department of Electrophysiology, 251 Pomorska Street, 92-213 Lodz, Poland

⁸Columbia University, Department of Biomedical Informatics, 622 W 168th St PH20 3720, New York, NY 10032, USA

⁹Brigham and Women's Hospital, Department of Cardiology, 75 Francis St, Boston MA 02115, USA

¹⁰Brigham and Women's Hospital, Department of Medicine, Division of General Internal Medicine, 75 Francis St, Boston MA 02115, USA

¹¹Harvard Chan School of Public Health, Department of Health Policy and Management, 677 Huntington Ave, Boston, MA 02115, USA

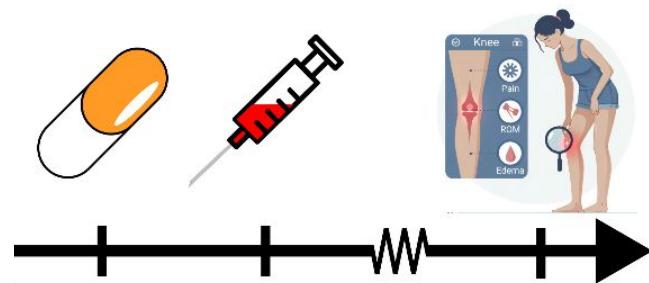
*Correspondence address. Arkadiusz Sitek, 55 Fruit St., Suite 427 Boston, MA, 02114 USA. E-mail: sarkadiu@gmail.com

Abstract

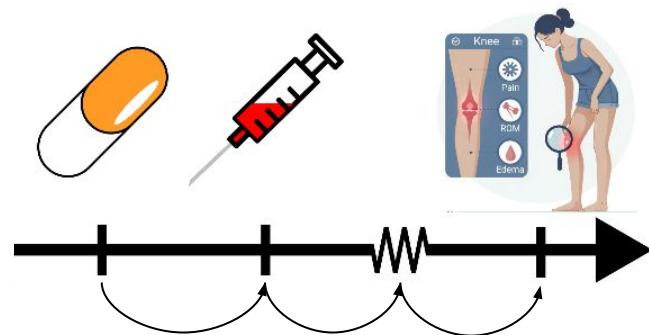
Background: Hospitals struggle to predict critical outcomes. Traditional early warning systems, like NEWS and MEWS, rely on static variables and fixed thresholds, limiting their adaptability, accuracy, and personalization.

Methods: We previously developed the Enhanced Transformer for Health Outcome Simulation (ETHOS), an artificial intelligence (AI) model that tokenizes patient health timelines (PHTs) from electronic health records and uses transformer-based architectures to predict future PHTs. ETHOS is a versatile framework for developing a wide range of applications. In this work, we develop the Adaptive

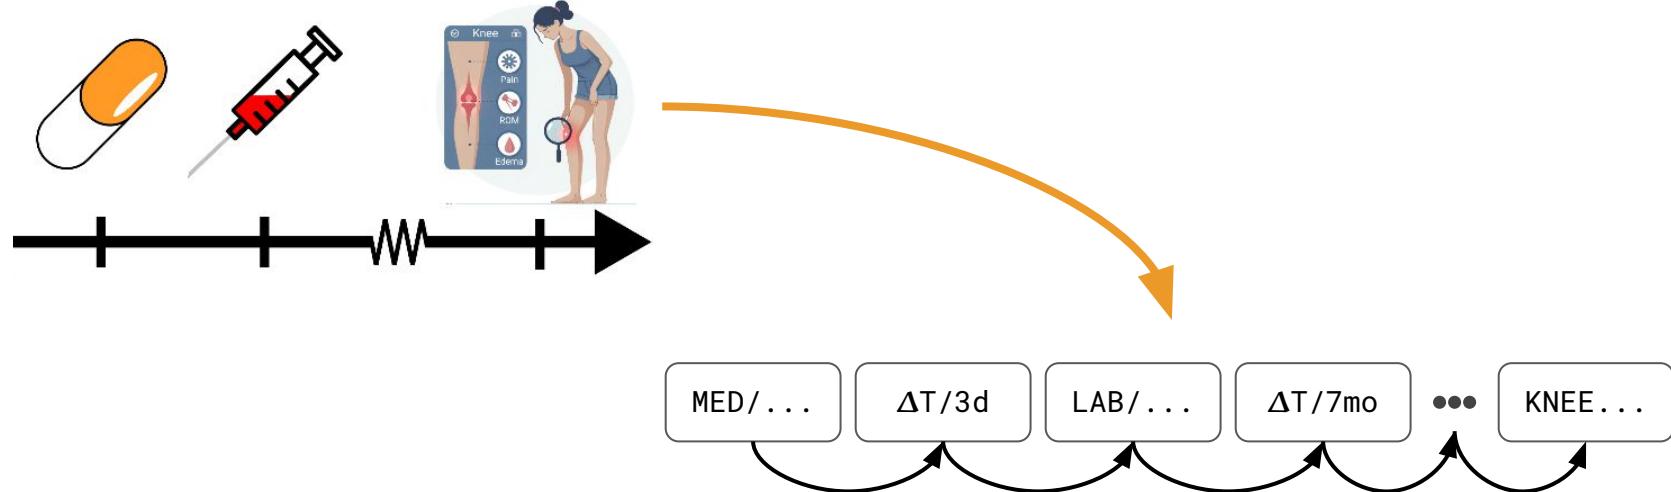
What observation comes next?



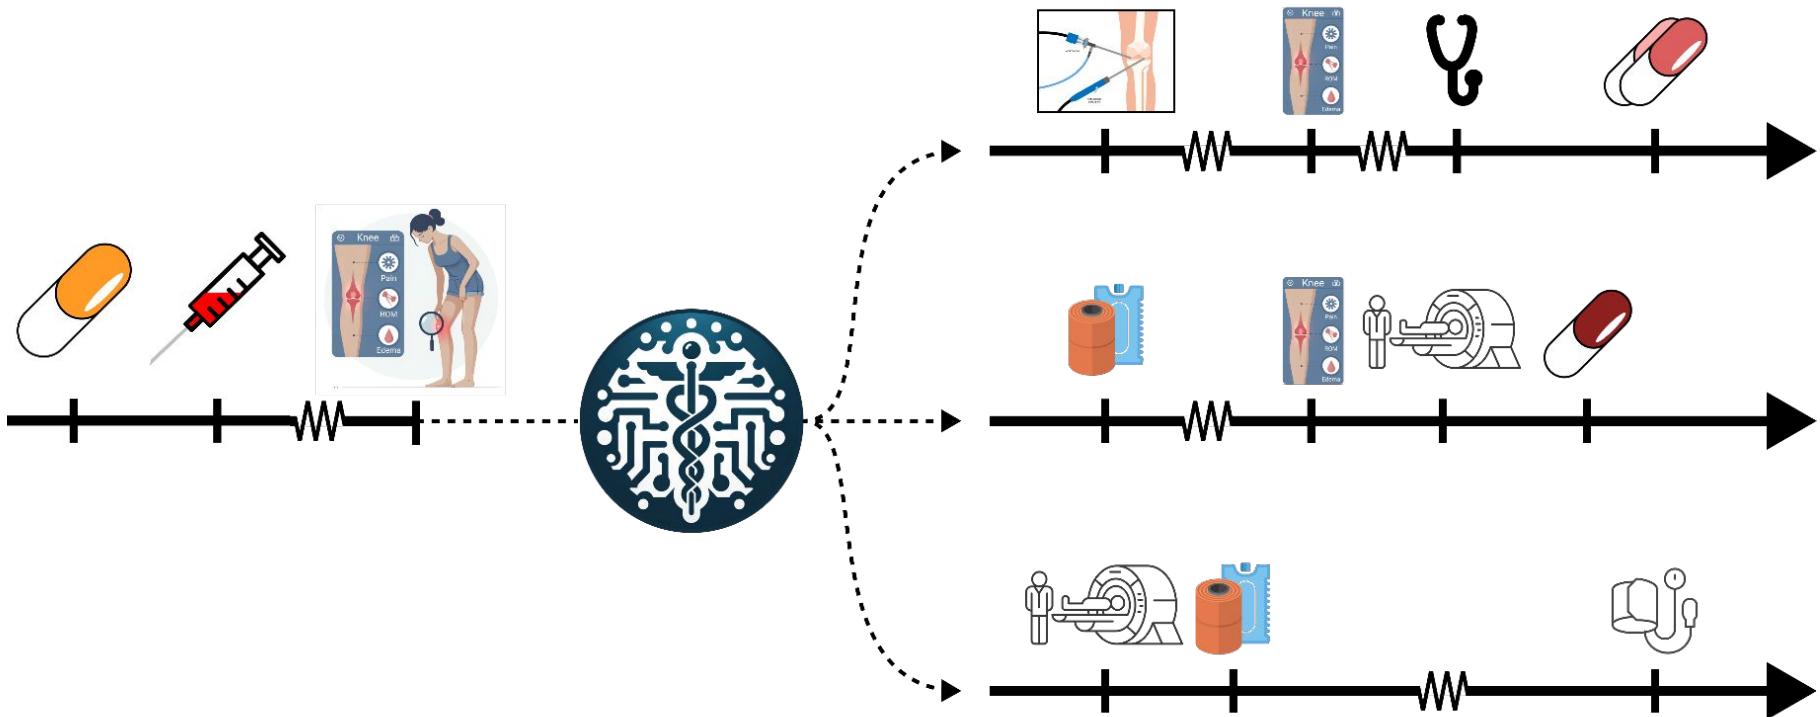
What observation comes next?



What observation comes next?



This is a foundation model because we can reason over synthetic futures:



An EHR Foundation Model Recipe

1. Convert to MEDS.
2. Use MEDS-EIC-AR.

The screenshot shows the GitHub repository page for `MEDS_EIC_AR`. The repository is public and has 10 branches and 13 tags. The main branch is the latest. The repository was created by `mmcdermott` and has 148 commits. The repository is described as a MEDS, "Everything-is-code" style Autoregressive Generative Model, capable of zero-shot inference. It includes sections for About, Releases, Packages, and Deployments. The README file is visible at the bottom.

About
A MEDS, "Everything-is-code" style Autoregressive Generative Model, capable of zero-shot inference.

Releases 13
0.1.10 (Latest)
2 weeks ago
+ 12 releases

Packages
No packages published
Publish your first package

Deployments 13
pypi 2 weeks ago
+ 12 deployments

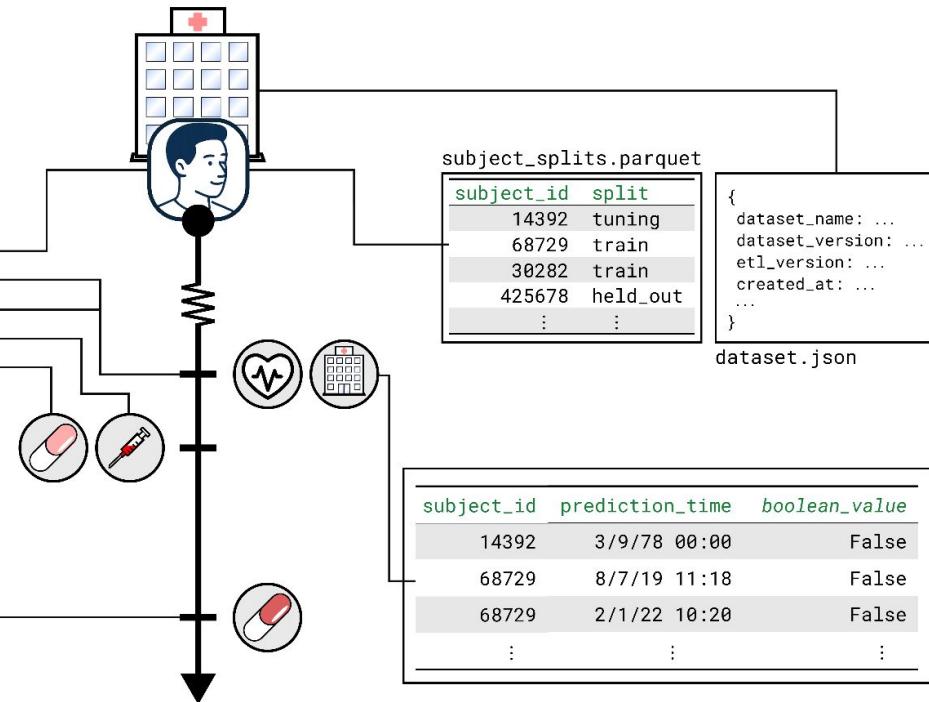
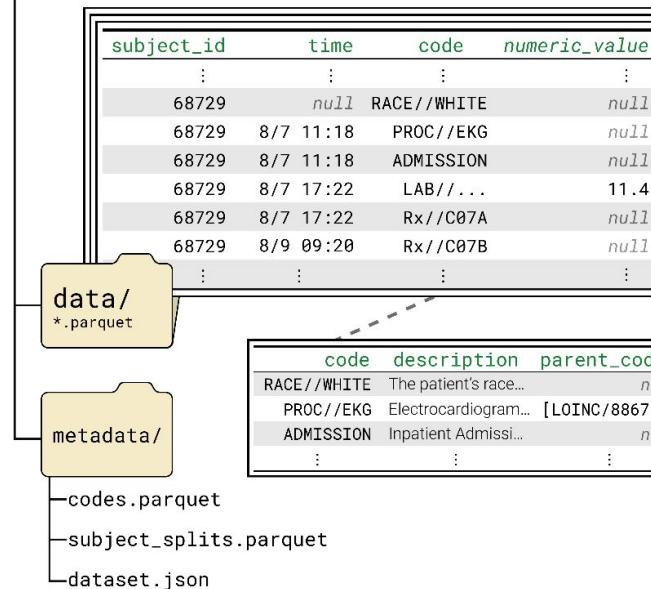
Languages
Python 100.0%

Installation

```
pip install MEDS-EIC-AR
```



MEDS



MEDS Reveals a Recipe for Autoregressive Foundation models on EHR Data

subject_id	time	code	numeric_value
:	:	:	:
68729	null	RACE//WHITE	null
68729	8/7 11:18	PROC//EKG	null
68729	8/7 11:18	ADMISSION	null
68729	8/7 17:22	LAB//...	11.4
68729	8/7 17:22	Rx//C07A	null
68729	8/9 09:20	Rx//C07B	null
:	:	:	:

Given a sequence of MEDS rows (ordered by time), can I predict the next row?

[main](#) [10 Branches](#) [13 Tags](#)[Go to file](#)[Add file](#)[Code](#)**About**

A MEDS, "Everything-is-code" style Autoregressive Generative Model, capable of zero-shot inference.

mmcdermott Merge pull request #63 from mmcdermott/codex/fix-issues-.... c36f1dd · 2 weeks ago 148 Commits

1. Install
pip install MEDS-EIC-AR

2. Data Pre-processing
MEICAR_process_data input_dir="\$RAW_MEDS_DIR" \
intermediate_dir="\$INTERMEDIATE_DIR" \
output_dir="\$FINAL_DATA_DIR"

3. Pre-train the model
MEICAR_pretrain datamodule.config.tensorized_cohort_dir="\$FINAL_DATA_DIR" \
output_dir="\$PRETRAINED_MODEL_DIR" \
datamodule.batch_size=32

4. Generate Trajectories
MEICAR_generate_trajectories \
output_dir="\$GENERATED_TRAJECTORIES_DIR" \
model_initialization_dir="\$PRETRAINED_MODEL_DIR" \
datamodule.config.tensorized_cohort_dir="\$FINAL_DATA_DIR" \
datamodule.config.task_labels_dir="\$TASK_ROOT_DIR/\$TASK_NAME" \
datamodule.batch_size=32

pip install MEDS-EIC-AR

